VARI: Robust Text-to-Video Generation via Group-Structured Intent Alignment

Mingxuan Zhu[1,2*†], Hohei Chan[2,3*], Shiyuan Liu[2],
Yuanxing Zhang[2], Lu Chen[1]

[1]Shanghai Jiao Tong University
[2]KlingAI
[3]South China University of Technology

arXiv:2610.18189v1 [cs.CV] 13 Oct 2026

Abstract

Modern text-to-video (T2V) models generate visually compelling videos, yet meaning-preserving changes in prompt wording can substantially alter semantic realization and overall quality. Existing prompt rewriting modifies the inference input, whereas conventional post-training constructs supervision from one prompt formulation at a time; neither explicitly optimizes reliability across equivalent forms. We introduce **VARI** (**V**erified **A**lignment over **R**epresentations of **I**ntent), a group-structured framework that treats automatically verified formulations of the same visual intent as a shared supervision unit. VARI builds matched form–seed rollout pools, cross-scores each video against all forms, and applies a wording-independent structured intent audit to identify valid realizations and hard semantic failures. Group-Consistent Distillation (GCD) aligns every verified form with interchangeable valid realizations, while Form-Contrastive Flow-DPO (FC-DPO) learns group-level preferences from audited positive–negative pairs, prioritizing same-seed, different-form contrasts. On HunyuanVideo-1.5, VARI attains the highest T2V-CompBench, VBench, and StoryEval scores among the evaluated methods. On prompt-form extensions of T2V-CompBench and VBench, it improves weakest-form alignment over the strongest prompt-level baselines by 2.45 and 3.55 points and reduces cross-form drop by 30.5% and 32.2%, respectively, while preserving visual quality. Human evaluation, natural-caption controls, prompt-enhancement experiments, diversity analyses, and transfer to LTX-2.3 further support the robustness and transferability of these gains, while VARI requires no additional inference-time module.

Introduction

Large text-to-video generators based on diffusion or flow matching now achieve strong visual fidelity, motion quality, and instruction following (Ho, Jain, and Abbeel 2020; Lipman et al. 2023; Kong et al. 2024). Yet robust semantic realization remains challenging: meaning-preserving changes in prompt wording can alter entity counts, spatial relations, action bindings, event realization, or overall generation quality. Such variation is distinct from desirable visual diversity, since appearance, composition, and motion details may change while the core visual intent should remain invariant. For example, consider "A boy reading in front of a bench" and "A bench is positioned behind a boy who is reading." They express the same spatial relation, yet a generator may realize it under one wording and violate it under the other.

Inference-time prompt rewriting can improve a selected input, but it does not train the generator itself to behave reliably across equivalent formulations. Reward- and preference-based post-training instead samples multiple rollouts for a prompt, scores them with automatic rewards, and learns from preferred and rejected examples (Liu et al. 2025b; Rafailov et al. 2023; Zhang et al. 2026; Liu et al. 2025a; Xue et al. 2025). However, supervision constructed from one formulation at a time cannot distinguish intent-level robustness from wording-specific success. Semantic-preserving prompt forms provide complementary linguistic evidence for constructing more reliable supervision.

We introduce **VARI** (**V**erified **A**lignment over **R**epresentations of **I**ntent), a group-structured intent-alignment framework that treats automatically verified prompt forms as a shared supervision unit. For each visual intent z, we construct $\mathcal{P}(z) = \{p_0, p_1, \ldots, p_K\}$, generate a matched form–seed rollout pool, and cross-score every candidate against all forms. A generation seed determines the initial sampling noise; reusing it across forms controls stochastic variation, so same-seed differences more directly expose wording sensitivity. Because equivalent forms share one structured semantic target, score variation for a fixed video is treated as *scoring-view dispersion*, rather than form-dependent ground truth. A frozen, wording-independent structured intent auditor checks atomic entities, counts, relations, action roles, and temporal commitments without observing prompt wording. The audit determines semantic type, while cross-form evidence ranks candidates within type, yielding reliable positives and visually plausible hard negatives.

As illustrated in Figure 1, *Group-Consistent Distillation* (GCD) trains every verified form on interchangeable valid realizations, whereas *Form-Contrastive Flow-DPO* (FC-DPO) optimizes audited success–failure preferences under every form, prioritizing same-seed, different-form contrasts. The resulting generator requires no additional inference-time module. On HunyuanVideo-1.5 (Wu et al. 2025a), VARI improves T2V-CompBench (Sun et al. 2025) by 1.64 points over the strongest prompt-level baseline and achieves the

*These authors contributed equally.
†This work was conducted during the author's internship at KlingAI.

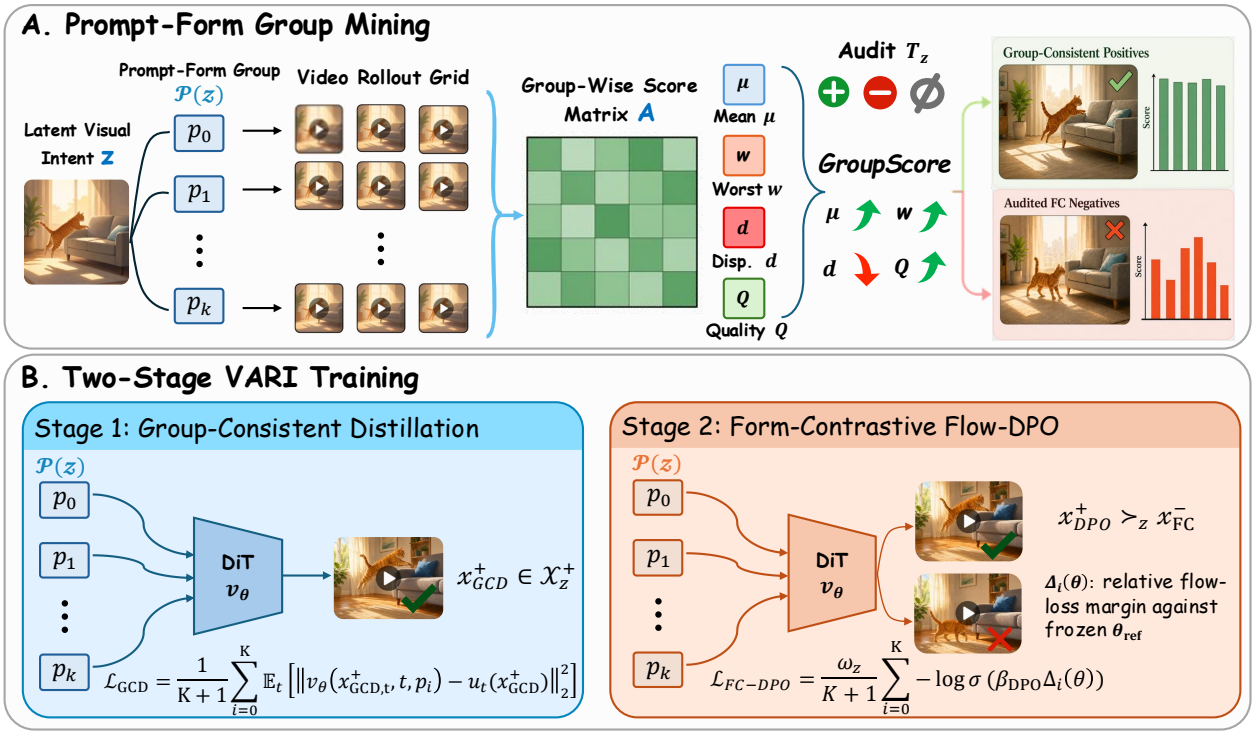

Figure 1: Overview of VARI. (A) Verified forms and shared seeds define a rollout grid. Cross-scoring provides mean support μ, weakest-view support w, dispersion d, and quality Q; a structured audit against T_z assigns semantic type, and GroupScore ranks candidates within type. (B) GCD distills retained positives across all forms, while FC-DPO lifts audited intent-level preferences to the full group, prioritizing same-seed, different-form pairs.

highest VBench (Huang et al. 2024) quality and StoryEval (Wang et al. 2025) scores among the evaluated methods. Prompt-form extensions of T2V-CompBench and VBench show weakest-form gains of 2.45 and 3.55 points, respectively. Human evaluation and held-out controls show that the gains persist under auditor replacement, natural captions, prompt enhancement, and transfer to LTX-2.3 while preserving seed diversity and sensitivity to non-equivalent prompt changes.

Our main contributions are threefold:

- **Verified prompt-form groups as supervision units.** We define prompt-form reliability over verified formulations of the same visually testable intent and use these groups, paired with shared seeds, as the supervision unit for T2V post-training.
- **Audited cross-form supervision.** We develop a compiler that cross-scores every rollout against all verified forms and uses a wording-independent structured intent audit to separate semantic validity from hardness ranking, yielding group-consistent positives and audited hard negatives.
- **Complementary group-level objectives.** We introduce GCD, which trains every verified form on interchangeable valid realizations, and FC-DPO, which optimizes audited success–failure preferences under every form while prioritizing same-seed, different-form comparisons.

Related Work

Diffusion, flow, and video generators. Modern visual generators build on denoising diffusion, latent diffusion, flow matching, rectified flows, diffusion transformers, and classifier-free guidance (Ho, Jain, and Abbeel 2020; Rombach et al. 2022; Lipman et al. 2023; Liu, Gong, and Liu 2023; Peebles and Xie 2022; Ho and Salimans 2022). Video generation has progressed from early pixel- and latent-space diffusion systems, cascaded spatial–temporal upsampling, and motion adaptation (Ho et al. 2022a,b; Singer et al. 2022; Blattmann et al. 2023a; Wang et al. 2023; Guo et al. 2023; Zhang et al. 2023; Kondratyuk et al. 2023; Bar-Tal et al. 2024) to large transformer- and flow-based backbones with stronger video latent spaces, including HunyuanVideo-1.5, LTX-2.3, and Wan2.2 (Blattmann et al. 2023b; Chen et al. 2024; Ma et al. 2024; Yang et al. 2024; Zheng et al. 2024; Polyak et al. 2024; Kong et al. 2024; Wu et al. 2025a; HaCohen et al. 2025; Lightricks 2026; Wan Team 2025a,b). VARI operates after pretraining and reorganizes supervision for an existing video DiT rather than modifying its backbone.

Preference and reward-based alignment. Generative-model alignment includes direct preference objectives, policy-gradient fine-tuning, reward backpropagation, and reward-model-free denoising preference learning (Rafailov et al. 2023; Black et al. 2023; Wallace et al. 2024; Yang et al. 2023). Image-generation studies have developed pref-

erence datasets, reward models, and input-side prompt optimization (Kirstain et al. 2023; Xu et al. 2023; Wu et al. 2023b; Hao et al. 2023). For video, VADER, VideoDPO, VideoReward, OnlineVPO, DenseDPO, DanceGRPO, and VisionReward improve reward learning, preference construction, temporal granularity, or on-policy optimization (Prabhudesai et al. 2024; Liu et al. 2025b,a; Zhang et al. 2026; Wu et al. 2025b; Xue et al. 2025; Xu et al. 2026). Recent LocalDPO, Diffusion-APO, and Dual-IPO further study localized supervision, trajectory-aware updates, and iterative reward–generator alignment (Huang et al. 2026; Zhu et al. 2026; Yang et al. 2026). VARI instead structures supervision across verified linguistic realizations of one visual intent. Main comparisons use VideoDPO and OnlineVPO because both admit matched HunyuanVideo-1.5 data and generation budgets; the newer methods change the underlying data or optimization pipeline.

The VARI Framework

Semantic-Preserving Prompt-Form Groups

Let z denote a visual intent and let p_0 be its original prompt. A verified prompt-form group

$$\mathcal{P}(z) = \{p_0, p_1, \dots, p_K\} \tag{1}$$

contains variants preserving the same structured target T_z, a set of atomic, visually testable commitments over objects, actions, attributes, relations, roles, and temporal order; unconstrained appearance may vary. Gemini 3.5 Flash generates ten training rewrites, Qwen3.6-27B verifies critical fields, and filtering retains four forms spanning lexical, syntactic, and caption-style variation; the roles are exchanged for benchmark groups. Human audit validates 190/200 training and 98/100 benchmark pairs, and 97/100 complete benchmark groups preserve T_z. Removing the three flagged groups changes headline robustness metrics by at most 0.11 points. Five forms and four shared seeds are adopted from held-out scale diagnostics, which show diminishing gains beyond five forms; details are in the supplement.

Each intent uses a seed set $\mathcal{S}_z$ shared across all verified forms, so form comparisons reuse the same initial sampling noises. The complete rollout pool is

$$\mathcal{X}(z) = \{G_{\theta_{\text{roll}}}(p_i, s) : p_i \in \mathcal{P}(z),\ s \in \mathcal{S}_z\}. \tag{2}$$

Here $G_{\theta_{\text{roll}}}$ is the frozen rollout checkpoint used to construct supervision, and s indexes the initial sampling noise. Mining uses the full within-intent form–seed pool; candidates are never paired across intents.

Cross-Form Score Geometry

Every candidate is evaluated against every verified form, yielding a score matrix $\mathbf{A}_z \in \mathbb{R}^{|\mathcal{X}(z)|\times(K+1)}$ with

$$[\mathbf{A}_z]_{x,i} = A_i(x) = \rho_A R_{\text{VQA}}(x, p_i) + (1-\rho_A) R_{\text{VS}}(x, p_i), \tag{3}$$

where VQAScore and VideoScore provide prompt-conditioned alignment signals (Lin et al. 2024; He et al. 2024). A separate low-weight quality term is

$$Q(x) = \rho_Q R_{\text{DOVER}}(x) + (1-\rho_Q) R_{\text{motion}}(x), \tag{4}$$

where R_{motion} summarizes motion smoothness, dynamics, flicker, and subject consistency (Wu et al. 2023a; Huang et al. 2024). Define the cross-form profile

$$\begin{aligned} \mu(x,z) &= \frac{1}{K+1}\sum_{i=0}^{K} A_i(x), \quad & w(x,z) &= \min_i A_i(x), \\ b(x,z) &= \max_i A_i(x), & d(x,z) &= b(x,z) - w(x,z), \end{aligned} \tag{5}$$

Thus μ measures average cross-form support, w the weakest scoring view, b the strongest, and d their range. We rank candidates with $S_{\text{grp}} = \lambda_\mu \mu + \lambda_w w - \lambda_d d + \lambda_q Q$, using $\rho_A = \rho_Q = 0.6$ and validation-fixed $(\lambda_\mu, \lambda_w, \lambda_d, \lambda_q) = (0.50, 0.25, 0.15, 0.10)$. Before the more expensive semantic audit, validation-fixed geometry–quality rules shortlist likely positives with high w, low d, and adequate Q, and likely hard negatives with high b, nontrivial d, and adequate Q. Only the deduplicated union is audited; the exact thresholds are in the supplement. This prefilter does not assign semantic labels: it only removes candidates that cannot satisfy the final positive or negative criteria.

For semantic typing, frozen Gemini 3.5 Flash (`gemini-3.5-flash`) (Google DeepMind 2026) receives 24 temporally ordered frames and one atomic commitment $c_{z,j}$, without prompt strings, generating-form identity, or mining scores. A stateless score call returns probabilities for `A` (satisfied), `B` (violated), and `C` (uncertain); a separate diagnostic call records evidence only and never affects supervision. Write $T_z = \{c_{z,j}\}_{j=1}^{J_z}$, let $\mathcal{K}_z \subseteq \{1, \dots, J_z\}$ index the required visually testable commitments, and define $s_j = P(\texttt{A})$ and $v_j = P(\texttt{B})$. We aggregate

$$\begin{aligned} C_{\min}(x,z) &= \min_{j\in\mathcal{K}_z} s_j, \quad C_{\text{avg}}(x,z) = \frac{1}{J_z}\sum_{j=1}^{J_z} s_j, \\ V_{\max}(x,z) &= \max_{j\in\mathcal{K}_z} v_j. \end{aligned} \tag{6}$$

A disjoint 600-video, three-annotator split fixes thresholds to achieve at least 90% human precision; uncertain candidates abstain. The endpoint, request schema, frame sampling, and thresholds are frozen before mining, and the auditor is excluded from evaluation, checkpoint selection, and GroupScore tuning; Qwen3.6-27B (Qwen Team 2026) is used only for held-out replacement. Exact request and safety settings are in the supplement. Since all forms share T_z, d is scoring-view dispersion, not form-dependent truth.

Group-Structured Supervision Synthesis

The shortlist supplies the necessary geometry and quality conditions, while the audit determines semantic type. With all components normalized to $[0,1]$, a stable positive must pass the positive shortlist and satisfy $C_{\min} \geq .70$ and $C_{\text{avg}} \geq .80$; a hard negative must pass the negative shortlist and satisfy $V_{\max} \geq .75$. Uncertain candidates are excluded. We denote the resulting eligible sets by $\mathcal{C}_z^+$ and $\mathcal{C}_z^{\text{FC}-}$, and retain $\mathcal{X}_z^+$ and $\mathcal{X}_z^{\text{FC}-}$ as their top m_+ and m_- candidates by GroupScore. Thus the audit fixes semantic type and GroupScore only ranks utility within type; for negatives, this favors visually plausible audit-confirmed failures rather than the

most degraded videos. Exact set definitions are in the supplement.

Each rollout retains provenance $x_{i,s} = G_{\theta_{\text{roll}}}(p_i, s)$. For a retained pair r, define its GroupScore gap as $g_{z,r} = S_{\text{grp}}(x^+_{z,r}, z) - S_{\text{grp}}(x^-_{z,r}, z)$ and its confidence weight as $\omega_{z,r} = \text{clip}(g_{z,r}, \omega_{\min}, \omega_{\max})$. The compiler prioritizes audited positive–negative pairs generated with the same seed under different forms and satisfying a validation-fixed margin, then fills any remaining per-intent budget with top-gap audited pairs from the same intent; their union is Π_z. Thus "form-contrastive" describes generating-form provenance, while correctness is always judged against the common target T_z. The compiler emits one shared-target record per retained positive and group-lifted preference records for pairs in Π_z; exact set definitions and thresholds are in the supplement.

Group-Consistent Distillation: Shared-Support Contraction

For a video flow model with conditional velocity predictor v_θ, let $\ell_{\text{FM}}(x^+, p_i, t; \theta)$ denote the standard squared flow-matching error between $v_\theta(x^+_t, t, p_i)$ and target $u_t(x^+)$. Here x^+_t is the corrupted or interpolated positive latent at sampled time t, and one record shares the latent and target across forms. Let q^+_z sample the retained set $\mathcal{X}^+_z$; GCD optimizes

$$\mathcal{L}_{\text{GCD}} = \mathbb{E}_{z,\, x^+ \sim q^+_z,\, t} \left[\frac{1}{K+1} \sum_{i=0}^{K} \ell_{\text{FM}}(x^+, p_i, t; \theta) \right]. \tag{7}$$

Different records sample different valid positives and noise, so GCD contracts around a valid support set rather than one canonical video. The supplement shows that the per-record objective decomposes into fitting the group-mean prediction plus an explicit pairwise disagreement penalty; diversity and counterfactual controls test that this contraction does not erase valid variation or prompt sensitivity.

Form-Contrastive Flow-DPO: Group-Lifted Margin Separation

For a mined pair and sampled time t, let $\Delta_i(\theta, t)$ be the policy-versus-reference change in the rejected-minus-preferred flow-loss margin under form p_i, where the reference is the frozen post-GCD checkpoint. VARI lifts the intent-level relation $x^+_{z,r} \succ_z x^-_{z,r}$ to every verified form and instantiates the preference update with a flow-DPO kernel:

$$\begin{aligned} \mathcal{L}_{\text{FC-DPO}} = \mathbb{E}_{z,r,t} \left[\frac{\omega_{z,r}}{K+1} \sum_{i=0}^{K} -\log \sigma(\beta_{\text{DPO}} \Delta_i(\theta, t)) \right] \\ + \lambda_{\text{aux}}\, \mathbb{E}_{z,r,t} \left[\frac{1}{K+1} \sum_{i=0}^{K} \ell_{\text{FM}}(x^+_{z,r}, p_i, t; \theta) \right]. \end{aligned} \tag{8}$$

We set $\beta_{\text{DPO}} = 0.3$ and $\lambda_{\text{aux}} = 0.05$. Within each branch, the policy and reference evaluate the same sampled time and corrupted input. The first term separates positives from audited hard negatives under every form; because the nonlinear preference loss is averaged after each form-wise margin is computed, an easy form cannot cancel an unresolved form before the logistic penalty, and smaller margins receive stronger updates. The auxiliary term preserves positive support, while matched seeds reduce stochastic confounding.

Algorithm 1 Audited Prompt-Form Supervision Compiler

Require: Intent z, target T_z, forms $\mathcal{P}(z)$, shared seeds $\mathcal{S}_z$, budgets m_+, m_-, margin τ_{pair}, clip bounds $\omega_{\min}, \omega_{\max}$
Ensure: GCD records $\mathcal{D}^z_{\text{GCD}}$ and FC-DPO records $\mathcal{D}^z_{\text{FC}}$
for each form $p_i \in \mathcal{P}(z)$ and seed $s \in \mathcal{S}_z$ **do**
  Generate $x_{i,s}$ and retain provenance (i, s)
end for
Cross-score the pool to obtain $\mathbf{A}_z$, Q, and $(\mu, w, b, d, S_{\text{grp}})$
Form geometry–quality shortlists $\mathcal{U}^+_z$ and $\mathcal{U}^-_z$ using validation-fixed conditions
for each $x \in \mathcal{U}^+_z \cup \mathcal{U}^-_z$ after deduplication **do**
  Audit atomic commitments in T_z and assign positive, negative, or abstained type
end for
Rank eligible candidates within type; retain $\mathcal{X}^+_z$ and $\mathcal{X}^{\text{FC}-}_z$
$\mathcal{D}^z_{\text{GCD}} \leftarrow \{(\mathcal{P}(z), x^+) : x^+ \in \mathcal{X}^+_z\}$; $\mathcal{D}^z_{\text{FC}} \leftarrow \varnothing$
for matched-seed pairs first, then top-gap audited fallback pairs **do**
  $g \leftarrow S_{\text{grp}}(x^+, z) - S_{\text{grp}}(x^-, z)$
  if $g \geq \tau_{\text{pair}}$ and the budget is not full **then**
    Add $(\mathcal{P}(z), x^+, x^-, \text{clip}(g, \omega_{\min}, \omega_{\max}))$ to $\mathcal{D}^z_{\text{FC}}$
  end if
end for
return $(\mathcal{D}^z_{\text{GCD}}, \mathcal{D}^z_{\text{FC}})$

Experiments

Experimental Setup

Backbone and training pool. We use HunyuanVideo-1.5 (Wu et al. 2025a) with a unified 480p, 16:9, 121-frame, 24-fps, 50-step protocol. The training pool contains 5,000 independently constructed source intents balanced across attributes, relations, actions, interactions, numeracy, and multi-event composition; each has four verified rewrites and four shared seeds, yielding 100,000 rollouts. All trainable methods use matched groups and generation budgets. Offline methods share the same pool, whereas OnlineVPO performs on-policy refresh under the same total budget. Main comparisons are Base, VideoDPO, OnlineVPO, GCD, and full VARI, with additional controls in the ablation.

Evaluation principle. Mining scores are not reused for reporting. Benchmark results are recomputed with the official VBench, T2V-CompBench, and StoryEval evaluator implementations. VBench and T2V-CompBench use the same 480p, 16:9, 121-frame, 24-fps, 50-step generation configuration for every system; StoryEval follows its official 10s setting. These extensions introduce no learned evaluator: they reuse the official per-video scores unchanged and add only verified prompt-form grouping and intent-level aggregation, yielding VBench-R and T2V-CompBench-R.

Prompt-form robustness protocol. For each of the 1,400 T2V-CompBench prompts, we construct five verified forms. VBench-R uses matched official semantic dimensions; each rewritten form inherits the source category and structured target, and only the generation-side wording changes. T2V-CompBench-R reuses the official score for the inherited category, whereas VBench-R reuses the matched official

Method	VBench Total	VBench Qual.	VBench Sem.	T2V-C Avg.	StoryEval
Base	80.11	82.83	69.23	47.69	32.11
VideoDPO	81.06	83.68	70.59	49.60	34.62
OnlineVPO	81.60	84.19	71.24	50.49	35.31
GCD	81.15	83.46	71.91	50.48	35.73
VARI	**82.67**	**84.91**	**73.71**	**52.13**	**37.58**

Table 1: Standard benchmark results. T2V-C averages seven official categories; StoryEval reports strict event completion.

semantic dimension and reports official VBench quality separately on the same clips. Mean, Worst, and max–min Drop are aggregated within each intent. FIFR counts mixed-success groups in which at least one form passes and another fails a validation-fixed, category- or dimension-specific threshold; it is interpreted jointly with Mean and Worst rather than as standalone generation quality.

Data independence and natural-caption controls. Source intents are split before rewriting, so every form of an intent remains in one partition. VBench, T2V-CompBench, StoryEval, VATEX, and MSR-VTT prompts are excluded from training, calibration, and checkpoint selection. Normalized exact matching, lexical retrieval, and frozen sentence-embedding retrieval followed by manual review find no shared visual intent; training and benchmark groups also use disjoint rewrite instructions and verifier instances. The supplement reports the full overlap audit and held-out VATEX (Wang et al. 2019) and MSR-VTT (Xu et al. 2016) controls scored by VisionReward (Xu et al. 2026). VATEX/MSR-VTT weakest-form scores improve from 61.92/59.24 for OnlineVPO to 65.09/62.34 for VARI.

Audit, coverage, and statistics. The protocol is frozen before mining, and only the union of the validation-fixed positive and negative shortlists is audited. Across 5,000 intents, 87.4% yield an audited positive, 64.6% an audited negative, and 61.4% a preference pair; 71.7% of retained pairs are same-seed, different-form. VideoDPO, OnlineVPO, GCD, and VARI use three matched training seeds; full deviations and source-intent paired tests are reported in the supplement.

Compute efficiency. All trainable methods share the 100,000-video rollout budget. VARI adds one-time, parallelizable offline cross-scoring, shortlist auditing, and a second optimization stage, giving modest overhead over VideoDPO and the same cost regime as OnlineVPO, while leaving per-video inference cost unchanged; details are in the supplement.

Standard Benchmark Results

Table 1 gives the headline benchmark view. VBench provides quality and semantic anchors, T2V-CompBench measures compositional prompt following, and StoryEval evaluates strict event completion (Wang et al. 2025). VARI gives the strongest overall scores. GCD is comparatively stronger on semantic consistency, while the full method also improves quality and StoryEval, indicating complementary stages without trading away fidelity or event completion.

Dimension-Wise Diagnostics

Table 2 expands T2V-CompBench into all seven categories. OnlineVPO remains strongest on dynamic attributes and motion binding, whereas VARI gains most on constant attributes, spatial relations, action binding, interaction, and numeracy. This pattern matches the entity, relation, role, and count commitments in the structured target rather than a uniform improvement in low-level motion generation.

Prompt-Form Robustness Extensions

Table 3 is the central prompt-form robustness evaluation. Unlike the mining stage, the scores in this table are official benchmark scores aggregated over verified prompt-form groups. Let $s_i^{(b)}(z)$ denote the official score of benchmark b for the video generated from the i-th prompt form of intent z. We report

$$\begin{aligned} \mathrm{Mean}(z) &= \frac{1}{|\mathcal{P}(z)|}\sum_i s_i^{(b)}(z), \\ \mathrm{Worst}(z) &= \min_i s_i^{(b)}(z), \\ \mathrm{Drop}(z) &= \max_i s_i^{(b)}(z) - \min_i s_i^{(b)}(z). \end{aligned} \tag{9}$$

Form-Induced Failure Rate (FIFR) measures mixed-success groups:

$$\begin{aligned} f_z^{(b)} &= \mathbf{1}\Big[\min_i s_i^{(b)}(z) < \tau_b \ \wedge \ \max_i s_i^{(b)}(z) \geq \tau_b\Big], \\ \mathrm{FIFR} &= \frac{100}{N}\sum_{z=1}^{N} f_z^{(b)}. \end{aligned} \tag{10}$$

Here τ_b is fixed on held-out validation data for each inherited category or matched semantic dimension. FIFR is interpreted jointly with Mean and Worst because a uniformly weak model can also have low mixed-success inconsistency; VBench quality is reported separately as a preservation guardrail.

Against the strongest prompt-level baseline for each endpoint, VARI raises T2V-CompBench-R/VBench-R Worst by 2.45/3.55 points, reduces Drop by 1.21/2.77, and reduces FIFR by 3.78/2.99. Relative to GCD, it further improves T2V-R Mean/Worst by 0.86/0.72 and lowers FIFR by 1.07, while T2V-R Drop increases modestly from 2.52 to 2.76. Together, the results suggest that GCD supplies the main cross-form contraction, whereas FC-DPO raises the lower tail

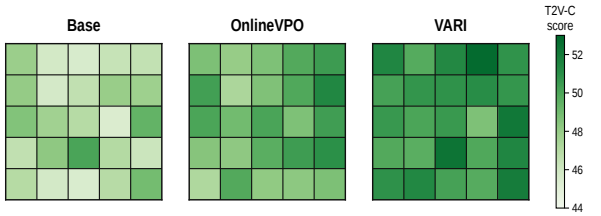

Figure 2: Representative mining-view score matrices for one intent group. Rows are generated candidates and columns are verified scoring forms. Cell variation is diagnostic evidence, not form-dependent ground truth for a fixed video.

Method	Cons. Attr.	Dyn. Attr.	Spatial Rel.	Motion Bind.	Action Bind.	Interaction	Numeracy	Avg.
Base	79.82	7.47	53.11	40.29	69.31	60.57	23.24	47.69
VideoDPO	81.34	8.25	55.56	41.89	71.24	63.49	25.42	49.60
OnlineVPO	82.31	**10.04**	56.72	**43.91**	72.13	64.49	23.81	50.49
GCD	83.26	8.66	56.97	39.88	72.71	66.61	25.29	50.48
VARI	**86.64**	9.58	**58.55**	40.42	**75.06**	**67.78**	**26.87**	**52.13**

Table 2: T2V-CompBench categories. VARI gives the best average; OnlineVPO leads dynamic attribute and motion binding.

Method	T2V-CompBench-R				VBench-R Semantic				VBench-R Quality
	Mean↑	Worst↑	Drop↓	FIFR↓	Mean↑	Worst↑	Drop↓	FIFR↓	Mean↑
Base	47.18	44.47	4.92	11.14	68.29	61.27	11.74	11.38	82.78
VideoDPO	49.54	47.08	3.97	8.93	70.16	65.18	9.19	8.72	83.51
OnlineVPO	50.09	46.95	4.12	8.21	70.79	65.07	8.61	8.26	83.94
GCD	50.39	48.81	**2.52**	5.50	71.36	67.12	6.51	6.08	83.47
VARI	**51.25**	**49.53**	2.76	**4.43**	**72.68**	**68.73**	**5.84**	**5.27**	**84.24**

Table 3: Prompt-form group extensions of the official evaluators; VBench quality is reported separately.

and removes additional mixed-success failures, with only a small trade-off in the average within-group range. Intent-level VBench-R gains are significant (Worst 95% CI $[2.66, 4.43]$, Holm-adjusted $p < 0.001$; FIFR McNemar $p < 0.01$ after correction).

Figure 2 shows a held-out cross-score matrix: OnlineVPO raises many cells but remains uneven, whereas VARI is higher and more stable. These scores provide mining evidence only; the structured audit assigns semantic type against the shared target.

Distribution Analysis

Figure 3 summarizes all VBench-R groups. VARI shifts worst-form scores upward and tightens Drop; simultaneous Mean and Worst gains suggest that the improvement is not explained solely by a few strong forms or uniformly weak contraction.

Compatibility with Prompt Enhancement

PE rewrites the input, whereas VARI changes the generator. We freeze the official HunyuanVideo-1.5 template (Tencent Hunyuan 2025a,b) for both methods and evaluate against the original forms. Under identical enhanced prompts, VARI improves T2V-C/R-Worst from 57.63/53.09 to 59.91/56.28 and reduces R-Drop/FIFR from 6.91/7.64 to 4.57/4.14, isolating improved generator behavior from rewrite strength.

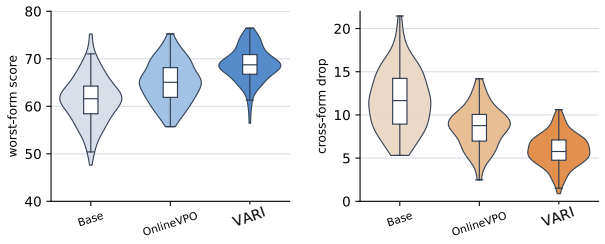

Figure 3: Group-level VBench-R semantic distributions. VARI raises worst-form scores and reduces cross-form drops.

Ablation Studies

Table 4 controls paraphrase exposure, worst-form optimization, target sharing, audit, hardness, and pair provenance. Exposure alone is insufficient; GCD provides the main contraction, while prompt-DPO after GCD raises Mean but degrades the lower tail. Audited group preferences improve Worst, Drop, and FIFR jointly, with matched-seed provenance adding the final gain.

Variant	Mean↑	Worst↑	Drop↓	FIFR↓
Reward-SFT	48.58	43.82	5.97	13.11
Multi-form prompt-DPO	50.08	47.03	4.34	8.61
GCD + prompt-DPO	50.82	48.64	3.28	6.19
Direct worst-form DPO	50.39	48.17	3.12	6.38
Form-specific distillation	49.93	47.51	3.69	7.19
Dispersion-only FC-DPO	50.69	48.35	3.58	6.91
GCD†	50.46	48.88	2.57	5.41
Intent-audit-only FC-DPO†	50.94	49.09	2.61	5.02
VARI w/o matched-seed pairs†	51.06	49.28	2.46	4.87
VARI (matched-seed prioritized)†	**51.17**	**49.51**	**2.29**	**4.44**

Table 4: Mechanism ablation on the fixed development split; † denotes three paired runs.

Supervision Validity and Boundary Controls

The blind audit in Table 5 separates semantic validity from dispersion: audited FC negatives reach 91/100 defect precision versus 58/100 for dispersion alone, showing that variation alone is not a reliable failure label. Auditor scores correlate with humans at $\rho = 0.63$; Qwen3.6-27B gives 0.73 retained-set Jaccard and 90/100 precision. VARI retains 96.5%/95.8% appearance/motion diversity and raises

Figure 4: Source/equivalent-form cases. Baselines violate spatial relation or action binding; VARI preserves both across forms. The panels show temporally ordered frame strips from the evaluated video clips.

the counterfactual margin from 11.36 to 13.90, providing evidence against collapse.

Mined type	Human precision	Visual usability	Independent agreement
Group-consistent positive	96/100	91/100	88/100
Dispersion-only negative	58/100	81/100	50/100
Audited FC negative	91/100	79/100	84/100

Table 5: Blind supervision audit; cells report valid counts out of $N = 100$ candidates per type.

Human Evaluation

We conduct a blinded pairwise comparison against OnlineVPO on 200 held-out intent groups. Three independent annotators compare anonymized video sets for intent faithfulness, cross-form consistency, visual quality, and overall preference; identities and side order are hidden, and ties are allowed. VARI obtains a 33.33-point overall win–loss margin (95% CI: 26.3–40.4; Krippendorff's $\alpha = 0.60$), with the largest advantage on consistency and a positive quality margin, matching the automatic Worst and Drop trends.

Cross-backbone transfer. On LTX-2.3 (Lightricks 2026), GCD raises T2V-C/R-Worst from 44.79/41.86 to 47.57/44.79 and lowers R-Drop from 5.48 to 3.51. Adding FC-DPO further reaches 48.29/45.76 and 3.01, while StoryEval rises to 33.37, reproducing the stage-wise trend on a second backbone.

Conclusion

We present VARI, a group-structured intent-alignment framework for robust text-to-video post-training. VARI treats verified formulations of the same visual intent as a shared supervision unit, separates semantic validity from hardness ranking through a wording-independent audit, and trains the generator with GCD and FC-DPO. The two stages play complementary roles: GCD shares interchangeable valid realizations across forms, while FC-DPO sharpens the preference boundary against audited semantic failures. Across HunyuanVideo-1.5 and LTX-2.3, VARI improves standard alignment and weakest-form reliability, reduces cross-form drop and mixed-success failures, and preserves visual quality, seed diversity, and sensitivity to genuine semantic changes. These results show that robustness to meaning-preserving prompt variation can be learned directly at the generator level, complementing inference-time prompt rewriting without adding deployment-time modules.

References

Bar-Tal, O.; et al. 2024. Lumiere: A Space-Time Diffusion Model for Video Generation. *arXiv preprint arXiv:2401.12945*.

Black, K.; Janner, M.; Du, Y.; Kostrikov, I.; and Levine, S. 2023. Training Diffusion Models with Reinforcement Learning. *arXiv preprint arXiv:2305.13301*.

Blattmann, A.; Rombach, R.; Ling, H.; Dockhorn, T.; Kim, S. W.; Fidler, S.; and Kreis, K. 2023a. Align Your Latents: High-Resolution Video Synthesis with Latent Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 22563–22575.

Blattmann, A.; et al. 2023b. Stable Video Diffusion: Scaling Latent Video Diffusion Models to Large Datasets. *arXiv preprint arXiv:2311.15127*.

Chen, H.; Zhang, Y.; Cun, X.; Xia, M.; Wang, X.; Weng, C.; and Shan, Y. 2024. VideoCrafter2: Overcoming Data Limitations for High-Quality Video Diffusion Models. *arXiv preprint arXiv:2401.09047*.

Google DeepMind. 2026. Gemini 3.5 Flash. https://ai.google.dev/gemini-api/docs/models/gemini-3.5-flash. Stable API model identifier: `gemini-3.5-flash`; accessed July 28, 2026.

Guo, Y.; et al. 2023. AnimateDiff: Animate Your Personalized Text-to-Image Diffusion Models without Specific Tuning. *arXiv preprint arXiv:2307.04725*.

HaCohen, Y.; Chiprut, N.; Brazowski, B.; Shalem, D.; Moshe, D.; Richardson, E.; Levin, E.; Shiran, G.; Zabari, N.; Gordon, O.; et al. 2025. LTX-Video: Realtime Video Latent Diffusion. *arXiv preprint arXiv:2501.00103*.

Hao, Y.; Chi, Z.; Dong, L.; and Wei, F. 2023. Optimizing Prompts for Text-to-Image Generation. In *Advances in Neural Information Processing Systems*.

He, X.; et al. 2024. VideoScore: Building Automatic Metrics to Simulate Fine-Grained Human Feedback for Video Generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*.

Ho, J.; Jain, A.; and Abbeel, P. 2020. Denoising Diffusion Probabilistic Models. In *Advances in Neural Information Processing Systems*.

Ho, J.; and Salimans, T. 2022. Classifier-Free Diffusion Guidance. *arXiv preprint arXiv:2207.12598*.

Ho, J.; Salimans, T.; Gritsenko, A.; Chan, W.; Norouzi, M.; and Fleet, D. J. 2022a. Video Diffusion Models. *arXiv preprint arXiv:2204.03458*.

Ho, J.; et al. 2022b. Imagen Video: High Definition Video Generation with Diffusion Models. *arXiv preprint arXiv:2210.02303*.

Huang, Z.; Zhang, K.; Ding, Y.; Gao, C.; Ding, R.; Chen, Y.; and Zuo, W. 2026. Mind the Generative Details: Direct Localized Detail Preference Optimization for Video Diffusion Models. *arXiv preprint arXiv:2601.04068*.

Huang, Z.; et al. 2024. VBench: Comprehensive Benchmark Suite for Video Generative Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Kirstain, Y.; Polyak, A.; Singer, U.; Matiana, S.; Penna, J.; and Levy, O. 2023. Pick-a-Pic: An Open Dataset of User Preferences for Text-to-Image Generation. *arXiv preprint arXiv:2305.01569*.

Kondratyuk, D.; et al. 2023. VideoPoet: A Large Language Model for Zero-Shot Video Generation. *arXiv preprint arXiv:2312.14125*.

Kong, W.; et al. 2024. HunyuanVideo: A Systematic Framework For Large Video Generative Models. *arXiv preprint arXiv:2412.03603*.

Lightricks. 2026. LTX-2.3: Official Inference and Training Repository. https://github.com/Lightricks/LTX-2. Checkpoint `ltx-2.3-22b-dev.safetensors`; accessed July 28, 2026.

Lin, Z.; Pathak, D.; Li, B.; Li, J.; Xia, X.; Neubig, G.; Zhang, P.; and Ramanan, D. 2024. Evaluating Text-to-Visual Generation with Image-to-Text Generation. *arXiv preprint arXiv:2404.01291*.

Lipman, Y.; Chen, R. T. Q.; Ben-Hamu, H.; Nickel, M.; and Le, M. 2023. Flow Matching for Generative Modeling. In *International Conference on Learning Representations*.

Liu, J.; Liu, G.; Liang, J.; Yuan, Z.; Liu, X.; Zheng, M.; Wu, X.; Wang, Q.; Xia, M.; Wang, X.; Liu, X.; Yang, F.; Wan, P.; Zhang, D.; Gai, K.; Yang, Y.; and Ouyang, W. 2025a. Improving Video Generation with Human Feedback. In *Advances in Neural Information Processing Systems*, volume 38.

Liu, R.; Wu, H.; Zheng, Z.; Wei, C.; He, Y.; Pi, R.; and Chen, Q. 2025b. VideoDPO: Omni-Preference Alignment for Video Diffusion Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Liu, X.; Gong, C.; and Liu, Q. 2023. Flow Straight and Fast: Learning to Generate and Transfer Data with Rectified Flow. In *International Conference on Learning Representations*.

Ma, X.; Wang, Y.; Jia, G.; Chen, X.; Liu, Z.; Li, Y.-F.; Chen, C.; and Qiao, Y. 2024. Latte: Latent Diffusion Transformer for Video Generation. *arXiv preprint arXiv:2401.03048*.

Peebles, W.; and Xie, S. 2022. Scalable Diffusion Models with Transformers. *arXiv preprint arXiv:2212.09748*.

Polyak, A.; et al. 2024. Movie Gen: A Cast of Media Foundation Models. *arXiv preprint arXiv:2410.13720*.

Prabhudesai, M.; Mendonca, R.; Qin, Z.; Fragkiadaki, K.; and Pathak, D. 2024. Video Diffusion Alignment via Reward Gradients. *arXiv preprint arXiv:2407.08737*.

Qwen Team. 2026. Qwen3.6-27B Model Card. https://huggingface.co/Qwen/Qwen3.6-27B. Unified multimodal checkpoint; accessed July 28, 2026.

Radford, A.; et al. 2021. Learning Transferable Visual Models From Natural Language Supervision. In *International Conference on Machine Learning*.

Rafailov, R.; Sharma, A.; Mitchell, E.; Ermon, S.; Manning, C. D.; and Finn, C. 2023. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*.

Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; and Ommer, B. 2022. High-Resolution Image Synthesis with Latent Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Singer, U.; et al. 2022. Make-A-Video: Text-to-Video Generation without Text-Video Data. *arXiv preprint arXiv:2209.14792*.

Sun, K.; Huang, K.; Liu, X.; Wu, Y.; Xu, Z.; Li, Z.; and Liu, X.

2025. T2V-CompBench: A Comprehensive Benchmark for Compositional Text-to-Video Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Tencent Hunyuan. 2025a. HunyuanVideo-1.5: Official Code and Prompt-Rewrite Utilities. https://github.com/Tencent-Hunyuan/HunyuanVideo-1.5. Accessed July 28, 2026.

Tencent Hunyuan. 2025b. HunyuanVideo 1.5 Prompt Handbook. https://github.com/Tencent-Hunyuan/HunyuanVideo-1.5/blob/main/assets/HunyuanVideo_1_5_Prompt_Handbook_EN.md. Accessed July 28, 2026.

Wallace, B.; Dang, M.; Rafailov, R.; Zhou, L.; Lou, A.; Purushwalkam, S.; Ermon, S.; Xiong, C.; Joty, S.; and Naik, N. 2024. Diffusion Model Alignment Using Direct Preference Optimization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Wan Team. 2025a. Wan: Open and Advanced Large-Scale Video Generative Models. *arXiv preprint arXiv:2503.20314*.

Wan Team. 2025b. Wan2.2: Official Code and Models. https://github.com/Wan-Video/Wan2.2. Accessed July 28, 2026.

Wang, J.; et al. 2023. ModelScope Text-to-Video Technical Report. *arXiv preprint arXiv:2308.06571*.

Wang, X.; Wu, J.; Chen, J.; Li, L.; Wang, Y.-F.; and Wang, W. Y. 2019. VATEX: A Large-Scale, High-Quality Multilingual Dataset for Video-and-Language Research. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*.

Wang, Y.; He, X.; Wang, K.; Ma, L.; Yang, J.; Wang, S.; Du, S. S.; and Shen, Y. 2025. Is Your World Simulator a Good Story Presenter? A Consecutive Events-Based Benchmark for Future Long Video Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Wu, B.; et al. 2025a. HunyuanVideo 1.5 Technical Report. *arXiv preprint arXiv:2511.18870*.

Wu, H.; Zhang, E.; Liao, L.; Chen, C.; Hou, J.; Wang, A.; Sun, W.; Yan, Q.; and Lin, W. 2023a. Exploring Video Quality Assessment on User Generated Contents from Aesthetic and Technical Perspectives. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 20144–20154.

Wu, X.; et al. 2023b. Human Preference Score v2: A Solid Benchmark for Evaluating Human Preferences of Text-to-Image Synthesis. *arXiv preprint arXiv:2306.09341*.

Wu, Z.; Kag, A.; Skorokhodov, I.; Menapace, W.; Mirzaei, A.; Gilitschenski, I.; Tulyakov, S.; and Siarohin, A. 2025b. DenseDPO: Fine-Grained Temporal Preference Optimization for Video Diffusion Models. *arXiv preprint arXiv:2506.03517*.

Xu, H.; Ghosh, G.; Huang, P.-Y.; Okhonko, D.; Aghajanyan, A.; Metze, F.; Zettlemoyer, L.; and Feichtenhofer, C. 2021. VideoCLIP: Contrastive Pre-training for Zero-shot Video-Text Understanding. In *Conference on Empirical Methods in Natural Language Processing*.

Xu, J.; Huang, Y.; Cheng, J.; Yang, Y.; Xu, J.; Wang, Y.; Duan, W.; Yang, S.; Jin, Q.; Li, S.; Teng, J.; Yang, Z.; Zheng, W.; Liu, X.; Zhang, D.; Ding, M.; Zhang, X.; Huang, S.; Gu, X.; Huang, M.; Tang, J.; and Dong, Y. 2026. VisionReward: Fine-Grained Multi-Dimensional Human Preference Learning for Image and Video Generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, 11269–11277.

Xu, J.; Liu, X.; Wu, Y.; Tong, Y.; Li, Q.; Ding, M.; Tang, J.; and Dong, Y. 2023. ImageReward: Learning and Evaluating Human Preferences for Text-to-Image Generation. In *Advances in Neural Information Processing Systems*.

Xu, J.; Mei, T.; Yao, T.; and Rui, Y. 2016. MSR-VTT: A Large Video Description Dataset for Bridging Video and Language. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.

Xue, Z.; Wu, J.; Gao, Y.; Kong, F.; Zhu, L.; Chen, M.; Liu, Z.; Liu, W.; Guo, Q.; Huang, W.; and Luo, P. 2025. DanceGRPO: Unleashing GRPO on Visual Generation. *arXiv preprint arXiv:2505.07818*.

Yang, K.; Tao, J.; Lyu, J.; Ge, C.; Chen, J.; Li, Q.; Shen, W.; Zhu, X.; and Li, X. 2023. Using Human Feedback to Fine-tune Diffusion Models without Any Reward Model. *arXiv preprint arXiv:2311.13231*.

Yang, X.; Yang, M.; Jia, G.; Qin, L.; Tan, Z.; and Li, H. 2026. Dual-IPO: Dual-Iterative Preference Optimization for Text-to-Video Generation. In *International Conference on Learning Representations*.

Yang, Z.; et al. 2024. CogVideoX: Text-to-Video Diffusion Models with an Expert Transformer. *arXiv preprint arXiv:2408.06072*.

Zhang, D. J.; Wu, J. Z.; Liu, J.-W.; Zhao, R.; Ran, L.; Gu, Y.; Gao, D.; and Shou, M. Z. 2023. Show-1: Marrying Pixel and Latent Diffusion Models for Text-to-Video Generation. *arXiv preprint arXiv:2309.15818*.

Zhang, J.; Wu, J.; Chen, W.; Ji, Y.; Xiao, X.; Huang, W.; and Han, K. 2026. Align Video Diffusion Model with Online Video-Centric Preference Optimization. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 6142–6152.

Zheng, Z.; Peng, X.; Yang, T.; Shen, C.; Li, S.; Liu, H.; Zhou, Y.; Li, T.; and You, Y. 2024. Open-Sora: Democratizing Efficient Video Production for All. *arXiv preprint arXiv:2412.20404*.

Zhu, J.; Chen, B.; Zhang, L.; Zhang, A.; Jiang, H.; and Huang, P. 2026. Diffusion-APO: Trajectory-Aware Direct Preference Alignment for Video Diffusion Transformers. *arXiv preprint arXiv:2605.07503*.

Supplementary Material

This supplement covers prompt-form construction and verification, audited supervision synthesis, objective analysis, implementation and data-independence protocols, statistical and alternative-explanation controls, extended results and ablations, human evaluation, the main-case video manifest, and four additional qualitative comparisons.

Prompt-Form Groups and Semantic Preservation

Visual intents and admissible forms

A visual intent z is represented by the unordered form set $\mathcal{P}(z) = \{p_0, p_1, \ldots, p_K\}$ and a structured target T_z. The target records entities and counts, attributes and states, actions and participant roles, interactions, spatial relations, temporal order, and explicit scene or camera constraints. The forms may differ lexically and syntactically but must preserve every field in T_z; the group is therefore an equivalence class of visually testable content, not loosely related captions.

Training source intents are instantiated from a deterministic, non-benchmark structured schema balanced across attributes, relations, actions, interactions, numeracy, and multi-event composition. Gemini 3.5 Flash (`gemini-3.5-flash`) generates ten candidates per source with a capped refill for malformed outputs. Qwen3.6-27B independently verifies every candidate against the source-derived T_z. For benchmark extensions, generator and verifier roles are exchanged and their prompts are disjoint from the training compiler. The endpoint IDs, access dates, and decoding settings are frozen in the released configuration.

Semantic-preserving prompt-form rewriting template

Input and goal. Given one T2V prompt, produce independently usable, linguistically diverse rewrites of the same visual event.
Preserve. Every entity and count, attribute/state, action and participant role, interaction, spatial relation, temporal order, scene constraint, and explicit camera/style/lighting/duration/text requirement.
Reject. Any added, removed, weakened, strengthened, disambiguated, or reversed visual commitment.
Permitted variation. Lexical substitution, clause reordering, role-preserving voice conversion, caption-like phrasing, and concise or expanded syntax without visual change.
Output. Return the source, each rewrite, and its transformation-family label in structured form.

The accepted transformation families are summarized in Table 6. They are used for audit and balance; the labels are not provided to the generator during training.

Family	Allowed surface variation
Lexical	Synonyms and local wording changes without semantic strengthening
Syntactic	Clause reordering, attachment changes, or role-preserving voice conversion
Caption-like	Natural descriptive-caption syntax with the same commitments
Concise	Removal of linguistic redundancy while retaining every visual constraint
Relation-focused	Reordering that foregrounds an existing spatial or interaction relation
Action-focused	Reordering that foregrounds an existing action or event sequence

Table 6: Prompt-form transformation families. A candidate is rejected whenever its visual content changes, regardless of its nominal family.

Structured equivalence verification

A rewrite is retained only when its relation to the source is unambiguous. Verification is field-wise rather than based on a single similarity score. This avoids accepting a fluent paraphrase that silently changes a count, swaps semantic roles, or modifies a spatial relation.

Visual-intent equivalence verification template

Task. Compare an original text-to-video prompt with one candidate rewrite. Judge semantic preservation only; do not judge writing style or predicted video quality.
Check independently. Entities and exact counts; attributes and states; actions; agent–patient roles; interactions; spatial relations; temporal order; scene constraints; explicit camera/style/lighting/duration/text constraints; and added, removed, or altered commitments.
Decision rule. Reject if any critical field differs, if a new visual commitment is introduced, if a required commitment disappears, or if a previously explicit relation becomes ambiguous. Treat an uncertain critical field as rejection rather than guessing.
Output structure. Return a status for every field, lists of added/removed/changed commitments, a final accept/reject decision, and a concise reason.

Illustrative audit example

For the source intent "A dog jumps over a low wooden fence in a grassy field," a rewrite may move the setting to the beginning or use a role-preserving paraphrase such as "In a grassy field, a dog clears a short wooden fence." A rewrite is rejected if it introduces a tracking camera, changes the fence height, adds a person, replaces *over* with *through*, or weakens the jump into generic movement. This example illustrates the field-wise verification policy.

Near-duplicate accepted forms are removed before group formation. Splits are created at the source-intent level before rewriting, so no forms of one intent can cross partitions. Benchmark-derived groups are excluded from training, threshold calibration, and checkpoint selection. Table 7 reports the compiler funnel. From 53,487 generated candidates, 2,713 malformed or near-duplicate outputs are removed and 34,318 pass field-wise verification; diversity-aware selection

retains four rewrites for each of 5,000 complete training groups.

Compiler stage	Count	Rate (%)
Generated candidates	53,487	100.00
Malformed / near-duplicate removed	2,713	5.07
Field-wise verifier inputs	50,774	94.93
Verifier-accepted candidates	34,318	64.16
Retained rewrites (four per source)	20,000	37.39
Complete five-form groups	5,000	100.00 intents

Table 7: Prompt-form generation and verification funnel. Candidate-stage rates use 53,487 as the denominator; complete-group coverage is reported over 5,000 source intents.

We additionally conduct a stratified human audit of 300 accepted source–rewrite pairs, with three independent annotators and majority vote. The sample contains 200 training pairs and 100 benchmark pairs. Table 8 reports the field-wise breakdown; Table 9 separates the two sources. A second audit samples 100 complete benchmark groups and asks whether all four rewrites preserve the common T_z. Ninety-seven groups pass in full. Excluding the three flagged groups changes every headline Worst, Drop, and FIFR result by at most 0.11 points, indicating that the benchmark conclusions are not driven by the audited failures.

Critical field	Equivalent / total	Precision (%)
Entity / count	59/60	98.33
Spatial relation	57/60	95.00
Action–role binding	58/60	96.67
Temporal order	56/60	93.33
Attribute / state / other	58/60	96.67
Overall	288/300	96.00

Table 8: Human semantic-equivalence audit of accepted prompt forms.

Audit subset	Equivalent / total	Precision (%)
Training accepted pairs	190/200	95.00
Benchmark accepted pairs	98/100	98.00
All accepted pairs	288/300	96.00
Complete benchmark groups (all four rewrites valid)	97/100	97.00

Table 9: Source-specific pair audit and complete-group benchmark audit. Group success requires all four rewrites to preserve the shared target.

Formal Objective Analysis

This section analyzes three properties of the VARI objectives: positive coverage under group-consistent mining, contraction induced by shared-target GCD, and the conservative form-wise preference surrogate optimized by FC-DPO.

Group-structured supervision compiler

Let $L = K + 1$ be the number of verified forms. The same intent-specific seed set $\mathcal{S}_z$ is reused across forms, and mining pools every candidate in $\mathcal{P}(z) \times \mathcal{S}_z$. For a candidate x, define its cross-form scoring profile

$$\begin{aligned} \mathbf{s}(x; z) &= \big(A(x, p_0), \ldots, A(x, p_K)\big), \\ \mu(x) &= L^{-1} \sum_i A_i(x), \qquad w(x) = \min_i A_i(x), \\ b(x) &= \max_i A_i(x), \qquad d(x) = b(x) - w(x). \end{aligned} \tag{11}$$

Surface-form-independent structured intent auditor

The primary auditor is Gemini 3.5 Flash (`gemini-3.5-flash`) (Google DeepMind 2026); `Qwen/Qwen3.6-27B` (Qwen Team 2026) is reserved for held-out replacement. Each 121-frame clip is represented by 24 uniformly sampled, temporally ordered RGB frames resized to 448×252. The structured target is deterministically serialized into atomic commitments $T_z = \{c_{z,1}, \ldots, c_{z,J_z}\}$ with field tags for entity/count, attribute/state, spatial relation, action–participant binding, interaction, temporal order, and explicit scene or camera constraints. Prompt strings, generating-form identity, and mining scores are withheld.

Score call. Each shortlisted candidate–commitment pair is evaluated by a fresh, stateless REST request whose output is restricted by a string-enum schema to one label token: `A` (satisfied), `B` (violated), or `C` (uncertain). We pin `thinkingLevel=minimal`, `candidateCount=1`, `maxOutputTokens=1`, `seed=202607`, `responseLogprobs=true`, and `logprobs=3`. In accordance with the Gemini 3.5 model guidance, `temperature`, `topP`, and `topK` are omitted rather than overridden; the stable model ID, complete request payload, and access date are archived with the code. Safety thresholds are `BLOCK_ONLY_HIGH` for the four adjustable categories: harassment, hate speech, sexually explicit content, and dangerous content. We renormalize the returned first-step probabilities over the token IDs for `A/B/C`. Let $s_j = P(\texttt{A})$ and $v_j = P(\texttt{B})$. Required visually testable fields form $\mathcal{K}_z$; we compute

$$\begin{aligned} C_{\min} &= \min_{j \in \mathcal{K}_z} s_j, \quad C_{\text{avg}} = J_z^{-1} \sum_j s_j, \\ V_{\max} &= \max_{j \in \mathcal{K}_z} v_j. \end{aligned} \tag{12}$$

A positive requires $C_{\min} \geq 0.70$ and $C_{\text{avg}} \geq 0.80$; a negative requires $V_{\max} \geq 0.75$. Other candidates abstain.

Atomic score query

You are given 24 temporally ordered frames and one atomic visual commitment. Judge only visible evidence. Return exactly one enum token: `A` if satisfied, `B` if violated, or `C` if uncertain. Do not provide any additional text.

Diagnostic call. A separate request may return a JSON object containing the categorical label, supporting frame indices, and a brief rationale. Diagnostic outputs support manual inspection and error analysis; they do not enter $C_{\min}$, C_{avg}, $V_{\max}$, threshold calibration, supervision construction, or checkpoint selection.

The thresholds are selected on a disjoint 600-video, three-annotator calibration split stratified by semantic field, maximizing retained coverage subject to at least 90% human precision for accepted positives and accepted semantic failures. Endpoint, frame sampling, label protocol, generation parameters, safety settings, field serialization, and thresholds are then frozen. The primary auditor is excluded from GroupScore-weight tuning, benchmark evaluation, and checkpoint selection. Qwen3.6-27B receives the identical frames and atomic commitments and is used only in the replacement analysis.

The alignment scores are normalized before aggregation. Because all forms share T_z, $d(x)$ is scoring-view dispersion rather than a form-dependent semantic label. The implemented GroupScore is

$$S_{\mathrm{grp}}(x,z) = \lambda_\mu \mu(x) + \lambda_w w(x) - \lambda_d d(x) + \lambda_q Q(x). \quad (13)$$

The non-audit parts of the positive and negative rules define

$$\begin{aligned} \mathcal{U}_z^+ &= \{x : w \geq .62,\ d \leq .18,\ Q \geq .55\}, \\ \mathcal{U}_z^- &= \{x : b \geq .58,\ d \geq .12,\ Q \geq .50\}. \end{aligned} \quad (14)$$

The auditor is invoked only on the deduplicated union $\mathcal{U}_z = \mathcal{U}_z^+ \cup \mathcal{U}_z^-$. This preserves the final candidate set: a rollout outside the corresponding shortlist cannot satisfy the complete positive or negative rule for any audit outcome.

VARI separates *semantic typing* from *hardness ranking*. With all components normalized to $[0,1]$, the implemented operators are

$$\Gamma^{+}(\phi_z(x)) = \mathbf{1}[x \in \mathcal{U}_z^+,\ C_{\min} \geq .70,\ C_{\mathrm{avg}} \geq .80], \quad (15)$$
$$\Gamma^{\mathrm{FC}-}(\phi_z(x)) = \mathbf{1}[x \in \mathcal{U}_z^-,\ V_{\max} \geq .75]. \quad (16)$$

The shortlist fixes necessary geometry and quality conditions, the audit defines semantic type, and GroupScore ranks candidates only within type. Among audit-confirmed failures satisfying $d \geq .12$, Top_{m_-} retains the highest-GroupScore candidates: visually plausible hard negatives rather than the most degraded or highest-dispersion failures. Every rollout stores its generating form and seed. The compiler consumes same-seed, different-form success–failure pairs first and fills the remaining budget with top-gap audited fallback pairs. "Form-contrastive" describes generating-form provenance, not form-dependent correctness of a fixed video. The clipped GroupScore gap becomes $\omega_{z,r}$.

Formally, let $g(x^+, x^-; z) = S_{\mathrm{grp}}(x^+, z) - S_{\mathrm{grp}}(x^-, z)$. The set Π_z^{match} contains same-seed, different-form pairs $(x_{j,s}^+, x_{i,s}^-)$ with $i \neq j$, $x_{j,s}^+ \in \mathcal{X}_z^+$, $x_{i,s}^- \in \mathcal{X}_z^{\mathrm{FC}-}$, and $g \geq \tau_{\mathrm{pair}}$. With top-gap fallback set $\Pi_z^{\mathrm{fallback}}$, the full set is $\Pi_z = \Pi_z^{\mathrm{match}} \cup \Pi_z^{\mathrm{fallback}}$; for each $r \in \Pi_z$, $g_{z,r} = g(x_{z,r}^+, x_{z,r}^-; z)$ and $\omega_{z,r} = \mathrm{clip}(g_{z,r}, \omega_{\min}, \omega_{\max})$. We use $\tau_{\mathrm{pair}} = 0.15$ and clip weights to $[0.05, 1.00]$.

Algorithm 2 Audited Prompt-Form Supervision Compiler

Require: Verified group $\mathcal{P}(z)$; pool $\mathcal{X}_z$; score matrix $\mathbf{A}_z$; limits m_+, m_-; pair margin m
Ensure: Shared-target records $\mathcal{D}_{\mathrm{GCD}}^z$ and preference records $\mathcal{D}_{\mathrm{FC}}^z$
Form $\mathcal{U}_z^+$ and $\mathcal{U}_z^-$ from the non-audit score/quality conditions
for each candidate $x \in \mathcal{U}_z^+ \cup \mathcal{U}_z^-$ **do**
    Audit T_z to obtain $(C_{\min}, C_{\mathrm{avg}}, V_{\max})$; form $\phi_z(x)$
end for
$\mathcal{C}_z^+ \leftarrow \{x \in \mathcal{U}_z^+ : C_{\min} \geq \tau_{\min}^+,\ C_{\mathrm{avg}} \geq \tau_{\mathrm{avg}}^+\}$
$\mathcal{C}_z^{\mathrm{FC}-} \leftarrow \{x \in \mathcal{U}_z^- : V_{\max} \geq \tau_{\mathrm{crit}}^-\}$
$\mathcal{X}_z^+ \leftarrow \mathrm{Top}_{m_+}(\mathcal{C}_z^+; S_{\mathrm{grp}})$
$\mathcal{X}_z^{\mathrm{FC}-} \leftarrow \mathrm{Top}_{m_-}(\mathcal{C}_z^{\mathrm{FC}-}; S_{\mathrm{grp}})$
$\mathcal{D}_{\mathrm{GCD}}^z \leftarrow \{(\mathcal{P}(z), x^+) : x^+ \in \mathcal{X}_z^+\}$
$\mathcal{D}_{\mathrm{FC}}^z \leftarrow \varnothing$
for each matched-seed pair, then each audited fallback pair **do**
    $g \leftarrow S_{\mathrm{grp}}(x^+, z) - S_{\mathrm{grp}}(x^-, z)$
    if $g \geq m$ **then**
        $\omega \leftarrow \mathrm{clip}(g, \omega_{\min}, \omega_{\max})$
        Add $(\mathcal{P}(z), x^+, x^-, \omega)$ to $\mathcal{D}_{\mathrm{FC}}^z$
    end if
end for
return $(\mathcal{D}_{\mathrm{GCD}}^z, \mathcal{D}_{\mathrm{FC}}^z)$

Compiler yield and pair provenance

The geometry/quality pre-screen sends only the union $\mathcal{U}_z^+ \cup \mathcal{U}_z^-$ to the structured auditor; candidates satisfying both branches are deduplicated before atomic queries. Across the 5,000 training intents, 4,371 provide at least one audited positive and 3,228 provide at least one audited negative. The compiler produces at least one preference pair for 3,069 intents; 1,302 additional intents contribute only to GCD. Among 5,907 retained preference pairs, 71.7% are same-seed, different-form pairs and 28.3% are audited fallback pairs.

Compiler outcome	Count	Rate (%)
At least one audited positive	4,371	87.42
At least one audited negative	3,228	64.56
At least one preference pair	3,069	61.38
GCD-only intent	1,302	26.04
No retained supervision	629	12.58
Matched-seed retained pairs	4,236	71.71
Fallback retained pairs	1,671	28.29

Table 10: Supervision yield. Intent rates use 5,000 source intents; pair-provenance rates use 5,907 retained preference pairs.

Commitments	N	Positive	Negative	Pair
$J_z = 3$–4	1,617	1,492 (92.3)	913 (56.5)	876 (54.2)
$J_z = 5$–6	2,356	2,068 (87.8)	1,539 (65.3)	1,474 (62.6)
$J_z \geq 7$	1,027	811 (79.0)	776 (75.6)	719 (70.0)
All intents	5,000	4,371 (87.4)	3,228 (64.6)	3,069 (61.4)

Table 11: Supervision yield by the number of atomic commitments. Cells report count (rate %). Min/max audit aggregation is deliberately more conservative for longer structured targets.

Proposition 1 (Coverage of an eligible positive)**.** *If a retained candidate x^+ satisfies $w(x^+) \geq \tau_w$ and $d(x^+) \leq \tau_d$, then*

$$\begin{aligned} A(x^+, p_i) &\geq \tau_w \quad \forall i, \\ |A(x^+, p_i) - A(x^+, p_j)| &\leq \tau_d \quad \forall i, j. \end{aligned} \tag{17}$$

Proof. Every score is at least the minimum w. All scores lie in $[w, b]$, whose diameter is $b - w = d$. □

This is the property required for target sharing: one favorable wording cannot by itself admit a positive. Mean-only ranking cannot provide this guarantee; for example, (μ, μ) and $(\mu - \delta, \mu + \delta)$ have the same mean but different worst scores and drops.

For an already audit-confirmed negative, a large-dispersion profile identifies a hard case rather than defining its semantic status. If $b(x^-) \geq \tau_b$ and $d(x^-) \geq \tau_d^-$, then there exist scoring forms p_h, p_ℓ such that

$$A(x^-, p_h) \geq \tau_b, \qquad A(x^-, p_\ell) \leq A(x^-, p_h) - \tau_d^-. \tag{18}$$

Thus the rejected set contains visually plausible, independently confirmed semantic failures with uneven scoring evidence, rather than universally poor videos or raw evaluator disagreement. The elementary bound

$$w(x) \geq \mu(x) - d(x) \tag{19}$$

also shows why improving the mean while reducing drop strengthens the lower tail.

Group-Consistent Distillation

Each intent retains a positive set $\mathcal{X}_z^+$, and each record samples $x^+ \sim q_z^+$ from that set. For time t and form p_i, let $h_i = v_\theta(x_t^+, t, p_i)$ and $u = u_t(x^+)$. With normalized weights $\alpha_i > 0$, the per-instance GCD loss is

$$\mathcal{L}_{\mathrm{GCD}}(x^+, z, t) = \sum_i \alpha_i \|h_i - u\|_2^2. \tag{20}$$

Proposition 2 (Shared-target decomposition)**.** *Let $\bar{h}_\alpha = \sum_i \alpha_i h_i$. Then*

$$\mathcal{L}_{\mathrm{GCD}} = \|\bar{h}_\alpha - u\|_2^2 + \sum_i \alpha_i \|h_i - \bar{h}_\alpha\|_2^2 \tag{21}$$

$$= \|\bar{h}_\alpha - u\|_2^2 + \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j \|h_i - h_j\|_2^2. \tag{22}$$

Proof. Write $h_i - u = (h_i - \bar{h}_\alpha) + (\bar{h}_\alpha - u)$ and expand the weighted squared norm. The cross term vanishes because $\sum_i \alpha_i (h_i - \bar{h}_\alpha) = 0$; the second equality is the weighted pairwise-variance identity. □

The first term fits the group-average prediction to the sampled valid target, while the second controls disagreement among verified forms. Sharing one target within a record therefore induces pairwise consistency, but sampling multiple retained positives across records does not define a unique canonical rendering. For uniform weights, $\partial \mathcal{L}_{\mathrm{GCD}} / \partial h_i = 2(h_i - u)/L$; a prediction-space step contracts each pairwise difference by $1 - 2\eta/L$ when $0 < \eta < L$.

Group-consistent mining is essential to this construction. If x^+ were selected only by its mean score, a low-scoring form could be hidden by favorable forms and target sharing would transfer a semantically unsuitable video to that form. Proposition 1 removes precisely this failure mode.

Form-Contrastive Flow-DPO

Let $\Delta_i(\theta)$ be the positive–negative flow-loss margin under form p_i, measured relative to the frozen post-GCD reference model. Define

$$\phi_\beta(\Delta) = -\log \sigma(\beta \Delta) = \log(1 + e^{-\beta \Delta}), \tag{23}$$

and, for a mined pair,

$$\mathcal{L}_{\mathrm{grp}}(z, x^+, x^-) = \omega_{z,r} \sum_i \alpha_i \phi_\beta(\Delta_i). \tag{24}$$

All experiments use $\alpha_i = 1/L$ and $\beta = \beta_{\mathrm{DPO}}$.

Proposition 3 (Conservative per-form surrogate)**.** *Because ϕ_β is convex,*

$$\phi_\beta\left(\sum_i \alpha_i \Delta_i\right) \leq \sum_i \alpha_i \phi_\beta(\Delta_i). \tag{25}$$

Proof. $\phi_\beta''(\Delta) = \beta^2 \sigma(\beta\Delta)\sigma(-\beta\Delta) > 0$, so Jensen's inequality applies. □

Averaging the nonlinear loss over forms is therefore stricter than applying one logistic loss to the average margin: an easy form cannot cancel a failed form before the nonlinearity. The form-wise gradient

$$\frac{\partial \mathcal{L}_{\mathrm{grp}}}{\partial \Delta_i} = -\omega_{z,r} \alpha_i \beta \sigma(-\beta \Delta_i) \tag{26}$$

has larger magnitude for smaller margins, so unresolved forms automatically receive stronger updates. The clipped pair weight $\omega_{z,r}$ reflects mining confidence while bounding the influence of unusually large GroupScore gaps.

Complementarity of the two stages

Two distinct controls

GCD controls conditional disagreement around positives sampled from a valid support set. It contains no rejected sample and does not establish a preference boundary. **FC-DPO** orders positives above independently audited negatives under every form; matched seeds reduce stochastic confounding. The stages therefore address complementary failure modes: shared positive support first, explicit rejection of verified hard failures second.

The analysis above characterizes the empirical objectives over the finite verified form group associated with each intent. Robustness to broader language variation is evaluated empirically through the natural-caption controls.

Computational Cost and Implementation Details

The complete supervision-construction and two-stage optimization flow is summarized in Algorithm 1 of the main paper. Here we provide the computational cost, frozen protocol settings, and implementation details needed for reproduction.

For M candidates and L prompt forms, cross-scoring requires ML prompt-conditioned evaluator calls. Computing the statistics is $O(ML)$, and sorting candidates is $O(M \log M)$ per intent. These quantities are computed during supervision construction and can be cached. The optimization stages use fixed targets and pairs, so their per-step complexity remains comparable to the corresponding flow-matching or flow-DPO update.

Implementation configuration

The transfer experiment uses the official `ltx-2.3-22b-dev.safetensors` checkpoint (Lightricks 2026).

Software environment. The Linux software stack comprised Python 3.11.11, PyTorch 2.10.0, Transformers 5.2.0, Diffusers 0.37.0, Accelerate 1.13.0, NumPy 1.26.4, and PyYAML 6.0.3.

Per-form mining alignment and quality are computed as

$$A_i(x) = 0.6R_{\mathrm{VQA}}(x, p_i) + 0.4R_{\mathrm{VS}}(x, p_i), \tag{27}$$

$$Q(x) = 0.6R_{\mathrm{DOVER}}(x) + 0.4R_{\mathrm{motion}}(x). \tag{28}$$

After score normalization, the current GroupScore coefficients are

$$(\lambda_\mu, \lambda_w, \lambda_d, \lambda_q) = (0.50, 0.25, 0.15, 0.10), \tag{29}$$

with the drop term entering negatively. These scores construct supervision; reported benchmark values are recomputed with the corresponding benchmark evaluators.

For a prompt form p_i and a mined pair (x^+, x^-), the flow-DPO margin is

$$\begin{aligned}\Delta_i(\theta) = {} & \ell_{\mathrm{FM}}(x^-, p_i; \theta) - \ell_{\mathrm{FM}}(x^+, p_i; \theta) \\ & - \ell_{\mathrm{FM}}(x^-, p_i; \theta_{\mathrm{ref}}) + \ell_{\mathrm{FM}}(x^+, p_i; \theta_{\mathrm{ref}}).\end{aligned} \tag{30}$$

where θ_{ref} is the frozen post-GCD checkpoint. The per-form update is

$$\begin{aligned}\ell_{\mathrm{FC}}(p_i, x^+, x^-) = {} & -\log \sigma(\beta_{\mathrm{DPO}} \Delta_i(\theta)) \\ & + \lambda_{\mathrm{aux}} \ell_{\mathrm{FM}}(x^+, p_i; \theta).\end{aligned} \tag{31}$$

with $\beta_{\mathrm{DPO}} = 0.3$ and $\lambda_{\mathrm{aux}} = 0.05$. The auxiliary term is chosen-only: it stabilizes the preferred support without directly fitting the rejected video.

For GCD, one sampled time t and one positive corruption are shared across all forms of a record, so the update isolates prompt conditioning while retaining a common target. For FC-DPO, the preferred and rejected branches share the same sampled time but use independent Gaussian corruptions. The policy and frozen reference evaluate the same corrupted input within each branch. This preserves a matched policy–reference comparison without coupling the preferred and rejected gradients through identical noise.

Protocol-level compute accounting

Table 13 compares the per-run execution pattern under the matched 100,000-video rollout budget used in the main experiments. Video generation is the dominant shared component. VideoDPO applies single-form scoring and one preference-training stage to the fixed pool; OnlineVPO interleaves on-policy generation, scoring, and policy/reference updates; VARI performs cross-form scoring once, audits only the union shortlist, and then runs GCD and FC-DPO offline. Because the stages have different execution patterns, we report matched workloads and execution structure rather than a single aggregate runtime metric. Repeated runs used for variance estimation are excluded from this per-run protocol comparison.

Under this matched budget, VARI incurs modest additional offline work over VideoDPO and remains in the same end-to-end cost regime as OnlineVPO. Its additional scoring and auditing are one-time and highly parallelizable, whereas OnlineVPO introduces sequential dependencies between generation, scoring, and policy updates. All three methods retain the same generator architecture and sampling configuration at inference, so VARI adds no per-video deployment cost.

Reproducibility and release. Headline trainable systems use three matched training seeds and fixed evaluation seeds. The four attribution-critical variants—GCD, intent-audit-only FC-DPO, VARI without matched-seed pairs, and full VARI—also use three paired runs from matched GCD branches; remaining ablations use one run on a shared development split unless stated otherwise. Compared methods share prompt groups, evaluation seeds, decoding settings, and evaluator versions whenever permitted by the method definition. VARI-specific settings are selected from held-out validation behavior and the accuracy–cost diagnostics reported below. The project implementation includes the VARI compiler, objectives, metrics, schemas, deterministic synthetic fixtures, tests, and manuscript-aggregate audits. External video backbones and evaluator implementations are connected through documented adapters and are not redistributed with this preprint. The synthetic fixtures support interface and regression tests, while the manuscript audits verify reported aggregate arithmetic. Source code, configurations, and evaluation scripts are not distributed with this preprint; any public release is subject to author and organizational approval. Any processed metadata release is also subject to applicable third-party licenses and organizational policies.

Evaluation Protocol

Intent-level robustness statistics

We extend the official T2V-CompBench and VBench evaluator implementations with verified prompt-form groups and intent-level aggregation, denoting the resulting protocols T2V-CompBench-R and VBench-R. Each verified rewrite inherits the source benchmark category and structured semantic target; the generator input is the only changed component. For a benchmark score $s_i^{(b)}(z)$ obtained for form i of

Item	Data and generation configuration
Main / transfer backbone	HunyuanVideo-1.5 / LTX-2.3 (22B dev)
Trainable component	Video DiT only; text encoders and VAE frozen
Training intents	5,000
Prompt-form groups	One source form and four verified semantic-preserving rewrites per intent
Generation seeds	Four intent-specific seeds reused across all verified forms
Initial rollout pool	100,000 videos (5,000 intents $\times$ 5 forms $\times$ 4 seeds)
Hunyuan benchmark generation	480p, 16:9, 121 frames at 24 fps, 50 sampling steps
Hunyuan qualitative generation	720p, 16:9, 121 frames at 24 fps, 50 sampling steps
LTX-2.3 transfer generation	480p, 16:9, 121 frames at 24 fps, 30 sampling steps
StoryEval generation	Official 10s HunyuanVideo evaluation setting

Item	Mining and optimization configuration
Retained positives / negatives	Up to two audited positives and two audited hard negatives per intent
Intent audit	Gemini 3.5 Flash enum score call; `Qwen/Qwen3.6-27B` replacement; frozen thresholds
Eligibility	$C_{\min}/C_{\text{avg}}/w/d/Q$ for positives; $V_{\max}/b/d/Q$ for negatives
Pairing	Margin 0.15; matched-seed pairs prioritized; weights clipped to $[0.05, 1.00]$
Initialization	GCD from step 2000; FC-DPO from fixed step-2400 GCD checkpoint
Updates / learning rates	400 / 5×10^{-6} (GCD); 600 / 2×10^{-6} (FC-DPO)
Optimizer / precision	AdamW, $(\beta_1, \beta_2) = (0.9, 0.95)$, $\epsilon = 10^{-8}$, wd 0.01; bf16
Batch / schedule	Global 48; no accumulation; clip 1.0; 50-step warmup then constant LR
DPO coefficients	$\beta_{\text{DPO}} = 0.3$; chosen-positive auxiliary coefficient 0.05
Checkpointing	Save every 300 updates; fixed terminal checkpoints, no test/qualitative selection
Trainable parameters	Full video DiT; text encoders and VAE frozen; no LoRA

Table 12: Core implementation configuration. Compared methods use matched generation settings and total sampling budgets whenever permitted by the method definition.

Method	Rollouts	Scoring / supervision	Optimization pattern	Inference overhead
VideoDPO	100k	Single-form reward and pair construction	Fixed-pool preference training	None
OnlineVPO	100k	Online reward and pair refresh	Interleaved generation and updates	None
VARI	100k	Cross-form scores; shortlist-only audit	Parallel offline GCD + FC-DPO	None

Table 13: Per-run protocol comparison under a matched video-rollout budget. API calls for the shortlist audit are logged separately by the compiler.

intent z, the robustness extension computes

$$\begin{aligned} \text{Mean}(z) &= L^{-1} \sum_i s_i^{(b)}(z), \\ \text{Worst}(z) &= \min_i s_i^{(b)}(z), \\ \text{Drop}(z) &= \max_i s_i^{(b)}(z) - \min_i s_i^{(b)}(z). \end{aligned} \tag{32}$$

Dataset-level Mean, Worst, and Drop are macro-averaged over visual intents. Form-Induced Failure Rate identifies groups for which at least one form passes a fixed benchmark-specific threshold while another form fails it:

$$f_z^{(b)} = \mathbf{1}\Big[\min_i s_i^{(b)}(z) < \tau_b \ \wedge \ \max_i s_i^{(b)}(z) \geq \tau_b\Big], \tag{33}$$

$$\text{FIFR} = 100\, N^{-1} \sum_{z=1}^{N} f_z^{(b)}. \tag{34}$$

The threshold is fixed independently of the compared method. T2V-CompBench-R uses a validation-frozen threshold for each inherited category, and VBench-R uses one for each matched semantic dimension. The reported FIFR is the macro-average of the resulting category- or dimension-wise failure rates. FIFR measures mixed-success inconsistency and can also be low when every form fails, so it is reported alongside Mean and Worst to retain absolute performance information. Standard benchmark scores and robustness statistics use the same generated samples and evaluator implementation within each comparison.

T2V-C category	Consist.	Dynamic	Spatial	Motion	Action	Interact.	Numeracy
τ_b (%)	78.07	7.18	50.22	36.31	65.09	56.37	20.19

Table 14: Validation-fixed success thresholds used by T2V-CompBench-R FIFR. VBench-R analogously uses the base-model validation median of each matched semantic dimension.

Comparison discipline

All methods receive the same prompt strings, evaluation seed sets, and generation configuration within each benchmark. The offline methods share a common mined rollout pool; OnlineVPO refreshes samples on-policy under the same total generation budget. Baseline checkpoint selection precedes qualitative inspection. Prompt-enhancement comparisons freeze each enhanced prompt and feed it identically to the Base and VARI-trained DiT. Quantitative VBench and T2V-CompBench results use the 480p, 121-frame, 24-fps

Hunyuan protocol with 50 sampling steps; the qualitative clips use 720p with the same frame count, frame rate, and 50-step sampler. StoryEval follows its official 10s HunyuanVideo evaluation protocol and serves as an auxiliary event-level diagnostic outside mining and checkpoint selection.

Statistical protocol and repeated runs

The statistical unit is the source intent, never an individual form or frame. Continuous paired differences use 10,000 intent-level bootstrap resamples for 95% confidence intervals and a two-sided Wilcoxon signed-rank test. FIFR uses paired bootstrap intervals for rate differences and McNemar's test. The eight primary robustness endpoints are corrected jointly with Holm's procedure. Human comparisons use cluster bootstrap over intent groups and retain ties.

Method	T2V-C Avg.	T2V-C-R Worst	VBench-R Worst
OnlineVPO	50.49 ± 0.23	46.95 ± 0.27	65.07 ± 0.29
GCD	50.48 ± 0.18	48.81 ± 0.21	67.12 ± 0.25
VARI	**52.13**± 0.19	**49.53**± 0.22	**68.73**± 0.26

Table 15: Mean $\pm$ standard deviation over three independent training runs.

Attribution-critical variant	**Mean**	**Worst**	**Drop**	**FIFR**
GCD	50.46 ± 0.18	48.88 ± 0.21	2.57 ± 0.13	5.41 ± 0.24
Intent-audit-only FC-DPO	50.94 ± 0.15	49.09 ± 0.19	2.61 ± 0.12	5.02 ± 0.22
VARI w/o matched-seed pairs	51.06 ± 0.19	49.28 ± 0.18	2.46 ± 0.16	4.87 ± 0.20
VARI (matched-seed prioritized)	**51.17**± 0.13	**49.51**± 0.16	**2.29**± 0.11	**4.44**± 0.18

Table 16: Mean $\pm$ standard deviation over three paired training branches on the fixed development split.

Relative to audit-only FC-DPO, full VARI changes Mean/Worst/Drop/FIFR by $+0.23/+0.42/-0.32/-0.58$. Relative to VARI without matched-seed pairs, the corresponding changes are $+0.11/+0.23/-0.17/-0.43$. The same ordering is retained across the three paired training branches in Table 16, supporting a modest final contribution from matched provenance and a larger contribution from audited group-level preference optimization.

Against the strongest prompt-level comparator for each endpoint, VARI improves T2V-CompBench-R/VBench-R Worst by 2.45/3.55 points, reduces Drop by 1.21/2.77 points, and reduces FIFR by 3.78/2.99 percentage points. For VBench-R, paired 95% intervals are $[2.66, 4.43]$ for Worst and $[-3.39, -2.14]$ for Drop; the Worst comparison has Holm-adjusted $p < 0.001$, and the FIFR reduction is significant under McNemar with adjusted $p < 0.01$. These paired VBench-R statistics are computed at the source-intent level from the final per-intent outputs rather than reconstructed from aggregate means.

Supervision Validity, Independence, and Boundary Tests

Blind audit of mined supervision

We sample 100 candidates from each mining type and hide method identity, generating form, scoring form, and all automatic scores. Three annotators judge the video only against T_z and separately record whether it is visually usable. Human-label precision denotes semantic validity for positives and semantic-defect precision for negatives. Independent agreement is the percentage on which a held-out audit model and the human majority assign the same target label.

Mined type	Human precision	Visual usability	Indep. agreement
Group-consistent positive	96/100	91/100	88/100
Low-score negative	76/100	47/100	68/100
Dispersion-only negative	58/100	**81/100**	50/100
Intent-audit negative	85/100	76/100	78/100
Audited FC negative	**91/100**	79/100	**84/100**

Table 17: Blind supervision audit. Bold marks the best result among negative-mining variants; each cell reports valid counts out of $N = 100$ candidates per type.

Across the 300 candidates used for correlation analysis, the primary Gemini auditor's $C_{\min}$ correlates with human correctness at Spearman $\rho = 0.63$. Let $\mathcal{N}_{\mathrm{G}}$ and $\mathcal{N}_{\mathrm{Q}}$ be the retained negative sets from Gemini 3.5 Flash and Qwen3.6-27B under the same per-intent budget $m_- = 2$. Their Jaccard overlap $|\mathcal{N}_{\mathrm{G}} \cap \mathcal{N}_{\mathrm{Q}}|/|\mathcal{N}_{\mathrm{G}} \cup \mathcal{N}_{\mathrm{Q}}|$ is 0.73; an independent review of 100 Qwen-retained negatives confirms 90 defects. Dispersion-only negatives remain substantially less precise despite high visual usability, supporting the separation between scoring-view dispersion and semantic typing.

Intent-disjoint construction and overlap audit

The 5,000 training intents are generated from a structured non-benchmark schema balanced across attributes, relations, actions, interactions, numeracy, and multi-event composition. Source-intent splits are frozen before rewriting. Evaluation groups use disjoint rewrite instructions and verifier instances. We normalize text for exact matching, retrieve lexical near-neighbors, and retrieve semantic neighbors using a frozen sentence encoder; every candidate above the validation-fixed threshold is manually reviewed against the structured intent.

Evaluation source	Retrieved candidates	Confirmed shared intents
VBench	13	0
T2V-CompBench	23	0
StoryEval	8	0
VATEX	17	0
MSR-VTT	26	0

Table 18: Training–evaluation overlap audit. Exact normalized overlap is zero for every source.

Seed diversity and non-equivalent counterfactuals

For each held-out intent and form, we generate eight seeds. Appearance features use the frozen CLIP ViT-L/14 image encoder (Radford et al. 2021): 16 uniformly sampled frames are independently encoded, ℓ_2-normalized, temporally averaged, and normalized again. Appearance diversity is the mean pairwise cosine distance across seeds. Motion features use the official VideoCLIP S3D checkpoint trained on HowTo100M (Xu et al. 2021): 32 ordered frames are encoded into the final

clip-level video embedding, normalized, and compared with cosine distance. Both diversity scores are averaged by intent and form and reported relative to Base.

For counterfactuals, we create 200 prompt pairs that change exactly one count, spatial relation, action role, or temporal order. Three annotators confirm that 195/200 initial pairs differ only in the intended commitment; the five failures are replaced before generation. The held-out VisionReward checkpoint (Xu et al. 2026), which is not used in mining or checkpoint selection, scores each video against both prompts. After mapping its score to the common 0–100 scale, the symmetric margin is $\frac{1}{2}[(R(x_p, p) - R(x_p, p')) + (R(x_{p'}, p') - R(x_{p'}, p))]$.

Method	Appearance div.	Motion div.	CF margin
Base	100.00	100.00	11.36
OnlineVPO	98.20	97.40	12.01
GCD	95.70	95.30	12.71
VARI	96.50	95.80	**13.90**

Table 19: Diversity retention relative to Base (%) and counterfactual semantic margin.

Change type	Base	OnlineVPO	GCD	VARI
Count	10.77	11.41	11.68	**13.12**
Spatial relation	13.26	14.01	14.82	**15.97**
Action-role swap	11.88	12.59	13.37	**14.53**
Temporal order	9.52	10.04	10.96	**11.98**

Table 20: Counterfactual margin by minimally changed semantic commitment.

VARI retains most seed-level appearance and motion diversity while increasing every counterfactual margin. These diagnostics provide evidence against mode contraction or indiscriminate suppression of prompt conditioning as the primary explanation for the robustness gains.

Expanded Quantitative Results

T2V-CompBench dimensions

Table 21 reports all seven displayed T2V-CompBench categories. The average is the arithmetic mean of the same seven entries.

Category-level observations. On T2V-CompBench, VARI obtains the highest overall average and leads five of the seven displayed categories: consistency, spatial relations, action, interaction, and numeracy. The remaining two categories expose a realistic trade-off rather than a uniformly monotone pattern. OnlineVPO is strongest on dynamic degree and motion binding, while GCD remains competitive on consistency, spatial relations, action, and interaction. The largest VARI gains occur in prompt-sensitive categories that require entity counts, relation preservation, or correct assignment of actions to participants.

Category	**Base**	**VideoDPO**	**OnlineVPO**	**GCD**	**VARI**
Consistency	79.82	81.34	82.31	<u>83.26</u>	**86.64**
Dynamic	7.47	8.25	**10.04**	8.66	<u>9.58</u>
Spatial	53.11	55.56	56.72	<u>56.97</u>	**58.55**
Motion	40.29	<u>41.89</u>	**43.91**	39.88	40.42
Action	69.31	71.24	72.13	<u>72.71</u>	**75.06**
Interaction	60.57	63.49	64.49	<u>66.61</u>	**67.78**
Numeracy	23.24	<u>25.42</u>	23.81	25.29	**26.87**
Average	47.69	49.60	<u>50.49</u>	50.48	**52.13**

Table 21: T2V-CompBench category scores under the unified evaluation protocol. Values are percentages; the table is transposed to preserve a stable single-column layout.

Ablations Linked to the Formal Analysis

Unless otherwise stated, the ablations in this section are measured on a fixed development split of held-out T2V-CompBench-R intent groups. Absolute values therefore need not exactly match the headline test-set results in the main paper, but every variant within an ablation table is evaluated on the same split and with the same evaluator.

Training objective

The objective ablation separates paraphrase exposure, shared support, semantic audit, and preference construction. Reward-SFT raises the mean but does not protect the weakest form. GCD produces the largest initial contraction. Adding ordinary prompt-DPO after GCD increases Mean but degrades Worst, Drop, and FIFR, whereas audited group-lifted FC-DPO improves all four endpoints over GCD, showing that Stage 2 contributes semantic boundary refinement rather than additional training alone.

Variant	Mean	Worst	Drop	FIFR
Reward-SFT	48.58	43.82	5.97	13.11
Prompt-DPO	49.91	46.77	4.49	9.24
FC-DPO	50.31	47.68	4.01	7.76
GCD	50.46	<u>48.88</u>	<u>2.57</u>	<u>5.41</u>
GCD + prompt-DPO	<u>50.82</u>	48.64	3.28	6.19
VARI	**51.17**	**49.51**	**2.29**	**4.44**

Table 22: Training-objective ablation on T2V-CompBench-R.

Causal alternatives

Table 23 matches data volume and evaluation split across simple alternatives. Multi-form prompt-DPO tests paraphrase exposure without group scoring; direct worst-form DPO tests explicit lower-tail optimization; form-specific and single-positive distillation test whether the gain is merely more SFT or one canonical target; inference reranking tests whether GroupScore helps without changing the generator; and the final rows isolate dispersion, audit, and matched-seed pairing.

Variant	Mean	Worst	Drop	FIFR
Reward-SFT	48.58	43.82	5.97	13.11
Prompt-level DPO	49.91	46.77	4.49	9.24
Multi-form prompt-DPO	50.08	47.03	4.34	8.61
Direct worst-form DPO	50.39	48.17	3.12	6.38
Form-specific distillation	49.93	47.51	3.69	7.19
Single-positive GCD	50.18	48.56	2.89	5.93
GCD	50.46	48.88	2.57	5.41
GroupScore reranking (no train)	49.69	46.45	4.13	8.69
Dispersion-only FC-DPO	50.69	48.35	3.58	6.91
Intent-audit-only FC-DPO	50.94	49.09	2.61	5.02
VARI w/o matched-seed pairs	51.06	49.28	2.46	4.87
VARI (matched-seed prioritized)	**51.17**	**49.51**	**2.29**	**4.44**

Table 23: Full causal ablation on the fixed T2V-CompBench-R development split.

The simple alternatives improve individual metrics but do not reproduce the final balance. Direct worst-form DPO contracts the lower tail but gives lower Mean and higher FIFR; reranking improves selection but not generator behavior; and dispersion-only negatives underperform audit-confirmed negatives. The paired repeats in Table 16 show that audited group-level preference optimization raises Worst and reduces mixed-success failures beyond GCD, while matched provenance adds a smaller complementary improvement in Worst, Drop, and FIFR.

Mining statistics

The component ablation matches Proposition 1 and the mean-only counterexample in the group-structured supervision analysis: mean-only selection can retain candidates with a weak tail, while adding w and d improves the lower-tail metrics. The quality term retains visually usable positives rather than the primary robustness signal; its main effect is to remove visibly weak positives while keeping robustness nearly unchanged.

Mining score	μ	w	$-d$	Worst↑	Drop↓	FIFR↓
Single-form	✗	✗	✗	41.08	16.89	14.94
Mean only	✓	✗	✗	44.17	13.43	11.18
Mean + worst	✓	✓	✗	46.12	10.21	8.24
Mean + worst - drop	✓	✓	✓	**46.95**	**7.48**	**5.96**
Full + Q	✓	✓	✓	46.87	7.64	6.08

Table 24: GroupScore-component ablation with T2V-CompBench-R evaluation.

Pair construction and group scale

Audited form-contrastive pairs are more informative than globally low-score negatives because they combine a genuine target violation with visual plausibility and uneven cross-form evidence. Increasing the number of verified forms improves coverage of wording variation but also increases cross-scoring cost. The gains largely saturate beyond five verified forms, so we adopt five forms and four seeds as the default accuracy–cost tradeoff.

Setting	Worst	Drop	FIFR
Random positive vs random negative	43.12	15.79	14.09
Best-group vs low-score	45.07	12.34	9.88
Best-group vs worst-group	45.74	9.83	7.41
Best-group vs form-contrastive	**46.88**	**7.54**	5.89
Top-m positives + mixed negatives	46.63	7.82	**5.74**
1 form, 4 seeds	41.23	16.21	14.98
3 forms, 4 seeds	44.93	10.71	9.37
5 forms, 2 seeds	45.19	9.97	8.71
5 forms, 4 seeds	46.81	7.73	6.07
7 forms, 4 seeds	**46.92**	**7.68**	**5.91**

Table 25: Pair-construction diagnostics (upper block) and rollout-scale diagnostics (lower block).

Prompt-form verification

Unverified rewrites may increase lexical diversity while silently changing the target intent. Structured verification improves the validity of cross-form supervision rather than merely increasing the number of prompts.

Group construction	Mean	Worst	Drop	FIFR
Single original form	49.91	41.24	16.23	14.98
Two verified forms	50.55	43.87	11.82	10.69
Five unverified rewrites	51.06	43.39	13.14	11.91
Five verified forms	**51.76**	**46.81**	**7.74**	**6.04**

Table 26: Effect of prompt-form verification and group size.

Prompt Enhancement and Cross-Backbone Evaluation

Official prompt-enhancement condition

HunyuanVideo-1.5 provides an official text-to-video prompt-rewriting instruction and a prompt-writing handbook for automatic enhancement (Tencent Hunyuan 2025a,b). The evaluation uses the official rewriting procedure without adding a method-specific semantic instruction. For each raw verified form, the enhanced text is frozen and supplied identically to the Base DiT and the VARI-trained DiT. Scoring remains anchored to the original benchmark intent, so the comparison tests whether generator-side alignment remains useful when inference-side enhancement is present.

Official PE evaluation contract

Input. The same raw prompt form for both compared generators.
Enhancement. Apply the official HunyuanVideo-1.5 text-to-video rewriting procedure once; do not manually edit one method's enhanced prompt or use enhancement during model selection.
Generation. Freeze the enhanced prompt and use the same inference configuration for the Base and VARI-trained DiT.
Evaluation. Score both outputs against the same original benchmark target. The enhanced text changes the generation-side conditioning, not the semantic reference used for comparison.

Input	Generator	T2V-C Avg.	R-Worst	R-Drop	R-FIFR
Raw	Base DiT	47.69	44.47	4.92	11.14
Raw	VARI DiT	52.13	49.53	**2.76**	4.43
Official PE	Base DiT	57.63	53.09	6.91	7.64
Official PE	VARI DiT	**59.91**	**56.28**	4.57	**4.14**

Table 27: Prompt-enhancement compatibility under a shared enhanced input.

Cross-backbone transfer and event-level diagnostic

The LTX-2.3 experiment (Lightricks 2026) evaluates whether the same two-stage training recipe transfers beyond the primary HunyuanVideo-1.5 backbone. Table 28 reports the full comparison under the matched LTX-2.3 protocol.

LTX-2.3	T2V-C	R-Worst	R-Drop	VBench-R	Story
Base	44.79	41.86	5.48	58.31	29.66
VideoDPO	46.31	43.19	4.41	60.08	31.02
OnlineVPO	46.84	43.04	4.82	60.67	31.48
+GCD	47.57	44.79	3.51	61.61	32.17
+VARI	**48.29**	**45.76**	**3.01**	**63.21**	**33.37**

Table 28: Cross-backbone validation on LTX-2.3. VBench-R reports semantic Worst; Story denotes StoryEval.

StoryEval uses the official 10s HunyuanVideo generation protocol and evaluator implementation. Its longer, event-centric prompts provide an auxiliary diagnostic of strict event completion outside supervision construction and checkpoint selection.

Method	Base	Reward-SFT	VideoDPO	OnlineVPO	VARI
StoryEval strict score	32.11	33.21	34.62	35.31	**37.58**

Table 29: StoryEval strict event-completion score.

Auxiliary Natural-Caption Controls

Verified VATEX multi-caption diagnostic

VATEX provides multiple English captions for the same source video, making it a useful auxiliary test bed for natural-caption grouping (Wang et al. 2019). We first collect candidate caption groups from the same video and then apply the same semantic-preservation verification used by VARI: retained captions must agree on the main entities, participant roles, principal action, core interaction or spatial relation, and temporal scope. Groups that describe different moments or omit a required relation are discarded. The verification retains 619 VATEX intent groups, all disjoint from the VARI training intents and checkpoint-selection data. Generation uses the same inference configuration as T2V-CompBench-R, while scoring uses held-out VisionReward only for this auxiliary diagnostic (Xu et al. 2026).

Method	Mean	Worst	Drop	FIFR
Base	63.77	59.66	6.18	9.21
VideoDPO	65.99	62.02	5.31	7.43
OnlineVPO	66.47	61.92	6.14	7.75
GCD	67.29	64.49	4.17	5.33
VARI	**68.16**	**65.09**	**4.03**	**4.85**

Table 30: Verified VATEX multi-caption diagnostic scored with held-out VisionReward.

Disjoint MSR-VTT control

MSR-VTT provides a stricter leakage-sensitive disjoint control (Xu et al. 2016). We remove duplicate video IDs, exact caption duplicates, and near-duplicate captions against the VARI training pool and benchmark-derived evaluation sets. The remaining captions undergo the same semantic-equivalence verification, retaining only groups that express the same visually testable intent. The filtering and verification retain 411 disjoint groups, which are evaluated with the same held-out reward protocol as VATEX.

Method	Mean	Worst	Drop	FIFR
Base	61.63	57.17	6.67	10.22
VideoDPO	63.61	59.33	6.07	8.76
OnlineVPO	64.05	59.24	6.37	9.00
GCD	65.21	61.69	4.96	6.33
VARI	**65.88**	**62.34**	**4.69**	**5.84**

Table 31: Disjoint MSR-VTT control under the same held-out reward protocol.

Human and Qualitative Evaluation Protocols

Human comparison

Human evaluation follows a blinded pairwise design on 200 held-out intent groups. Annotators receive the visual intent, the verified prompt forms, and anonymized video sets from two methods. Method identity and left/right placement are hidden and randomized. Separate questions assess intent faithfulness, consistency across equivalent forms, visual quality, and overall preference. Each comparison is judged by three independent annotators; ties are retained rather than forced into a directional choice. Win, tie, and loss are raw percentages over all judgments and therefore sum to 100%. Decisive-win rates excluding ties are reported only as an auxiliary statistic.

Human-evaluation questions

Intent faithfulness. Which set more accurately depicts the required entities, actions, attributes, and relations?
Prompt-form consistency. Which set preserves the same intended content more consistently across the equivalent forms?
Visual quality. Which set has better perceptual quality, temporal coherence, and fewer visible artifacts?
Overall preference. Considering faithfulness, consistency, and quality together, which set is preferred?

For the comparison with OnlineVPO, Table 32 reports integer outcomes and the percentages derived from them. The decisive-win rates excluding ties are 69.84% for intent

faithfulness, 73.11% for cross-form consistency, 63.93% for visual quality, and 70.24% overall. Confidence intervals use intent-cluster bootstrap; Krippendorff's α is computed on the three-way win/tie/loss labels.

Dimension	Counts			Raw shares (%)			W–L	95% CI	α
	W	T	L	W	T	L	(pt.)		
Intent faithfulness	352	96	152	58.67	16.00	25.33	33.33	[25.9, 40.6]	0.59
Cross-form consistency	378	83	139	63.00	13.83	23.17	39.83	[32.4, 47.1]	0.63
Visual quality	319	101	180	53.17	16.83	30.00	23.17	[14.9, 31.3]	0.57
Overall preference	347	106	147	57.83	17.67	24.50	33.33	[26.3, 40.4]	0.60

Table 32: Detailed blinded human evaluation: VARI versus OnlineVPO. Each dimension contains 600 judgments.

Qualitative cases and video comparison protocol

The main paper reports two held-out T2V-CompBench source prompts, one from spatial relationships and one from action binding. Each source is paired with one verified semantic-preserving form. The left/right prompt pairs are:

- **Spatial relation:** "A boy reading in front of a bench." and "A bench is positioned behind a boy who is reading."
- **Action binding:** "A man is reading a newspaper on a park bench, while a squirrel is climbing a tree." and "While a squirrel climbs a tree, a man reads a newspaper on a park bench."

All compared methods use the same prompt string and inference configuration within each row. The complete qualitative clips are generated with HunyuanVideo-1.5 at 720p, 121 frames, 24 fps, and 50 sampling steps. Quantitative VBench and T2V-CompBench tables use the corresponding 480p, 50-step protocol. The main-paper three-frame strips are sampled in temporal order from these complete clips. The twelve full clips underlying the frame-strip comparisons follow the organization summarized below; the video files are not distributed with this preprint.

Qualitative clip organization

The evaluated clips were organized as follows:

```
spatial/
  source/{base,videodpo,vari}.mp4
  equivalent/{base,videodpo,vari}.mp4
action_binding/
  source/{base,videodpo,vari}.mp4
  equivalent/{base,videodpo,vari}.mp4
```

All clips use the compared methods and the shared 720p generation protocol described above.

Spatial-relation case. Across both forms, Base and VideoDPO place the reader on the bench, replacing the requested *in-front-of/behind* relation with a frequent reading-on-a-bench configuration. One baseline clip additionally duplicates the reader, but the shared failure is the relation change rather than the mere presence of a counting artifact. The two VARI clips preserve the intended relation throughout the sequence while allowing an unconstrained posture change: the boy sits on the grass in one form and stands in front of the bench in the other. The complete clips also expose temporal behavior hidden by the three-frame layout. Baseline book pages can flip repeatedly or change geometry abruptly, whereas the VARI outputs maintain a more coherent reading interaction; the seated result contains a brief localized page turn and the standing result remains stable without changing the spatial arrangement.

Action-binding case. The baseline videos display several ways in which a partially correct scene can still violate the compositional intent. The man–newspaper event is often present, but squirrels may be duplicated, appear on the bench rather than the tree, or fail to maintain a single upward trajectory; one clip instead duplicates the reader. The two VARI clips preserve both assignments over time: the man continues reading and the squirrel remains on the tree while moving upward. The full clips also show stable newspaper handling with limited page motion, rather than a still-frame coincidence, and make the climbing progression visible from low to higher positions on the trunk.

Checkpoint selection precedes qualitative inspection. We then select held-out groups with clear weakest-form improvements and interpretable failure modes; the quantitative conclusions are based on the full benchmark results.

Scope of the public preprint

This preprint includes the full technical appendix and static qualitative panels. The twelve complete example clips and the project code/data are not distributed with it. The project implementation includes VARI schemas, the supervision compiler, GCD and FC-DPO objectives, robustness aggregation, prompt templates, deterministic synthetic fixtures, unit and smoke tests, a paper-to-code mapping, and manuscript-aggregate arithmetic audits. Documented adapters connect to external backbones, evaluator repositories, and pretrained weights.

Comparison	Forms	Material shown
Main-paper spatial relation	2×3	Displayed strips (6 evaluated clips)
Main-paper action binding	2×3	Displayed strips (6 evaluated clips)
Appendix paired-form cases	2×2	Static temporally ordered panels
Appendix dynamics cases	2×1	Static temporally ordered panels

Table 33: Qualitative evidence shown in this preprint. Each count is prompt forms $\times$ compared methods or intents, as applicable.

Additional qualitative results

Figure 5 provides four additional held-out comparisons using temporally ordered frames. Cases (a) and (b) compare the source prompt with one verified semantic-preserving form, while cases (c) and (d) focus on motion realization and temporal state change from the source prompt. These four cases are presented as static appendix panels; the twelve full clips underlying the main-paper frame strips are not distributed with this preprint.

Numeracy and group action. In panel (a), the numeral is a binding constraint rather than a stylistic preference.

Figure 5: Four additional held-out cases with ordered frames. (a) VARI preserves five salient robots and the group dance across both forms. (b) It maintains one coherent game layout and stable participants. (c) It moves the car past a stable church. (d) It realizes progressive snowman collapse and volume loss.

Base produces crowded groups with additional background robots, while VideoDPO commonly exposes only four dominant robots or changes the apparent count across frames and forms. VARI keeps five clearly identifiable robots in both prompt formulations and preserves the requested collective dancing behavior with stronger temporal and appearance consistency.

Multi-person interaction and object coherence. Panel (b) stresses scene-level logic rather than a single highlighted token. Base and VideoDPO frequently turn the singular board game into multiple disconnected, duplicated, or fragmented playing surfaces, weakening both object identity and the social interaction. VARI produces the most coherent table geometry and game layout, keeps the participants engaged with the same central activity, and yields cleaner faces, hands, and cafe structure across the sequence.

Qualitative audit cues

The four appendix cases are read against explicit visual commitments: exact entity count and shared group action in panel (a); stable participants and one coherent game object in panel (b); object displacement relative to fixed landmarks in panel (c); and a sustained, monotonic state transition in panel (d). The temporally ordered frames support these checks without substituting for the complete clips used in the main-paper cases.

Motion realization and scene preservation. For panel (c) in Figure 5, the decisive requirement is object motion relative to fixed landmarks. The Base strip provides little evidence of the car translating through the scene and is dominated by viewpoint change. VideoDPO keeps a visible car but shows weaker frame-to-frame appearance, scale, and trajectory continuity. VARI presents a clear progression from a distant car to a close pass while retaining the old church as a stable spatial reference.

Temporal state change. Panel (d) in Figure 5 tests whether the model executes a sustained transformation rather than merely rendering the correct initial object. The Base snowman remains essentially intact, and VideoDPO shows at most a small local dripping cue without a substantial change of state. VARI produces a pronounced, monotonic melting sequence: the body loses volume, the structure collapses, and the final frame visibly differs from the initial snowman.